

Virtual Memory

Kevin Webb

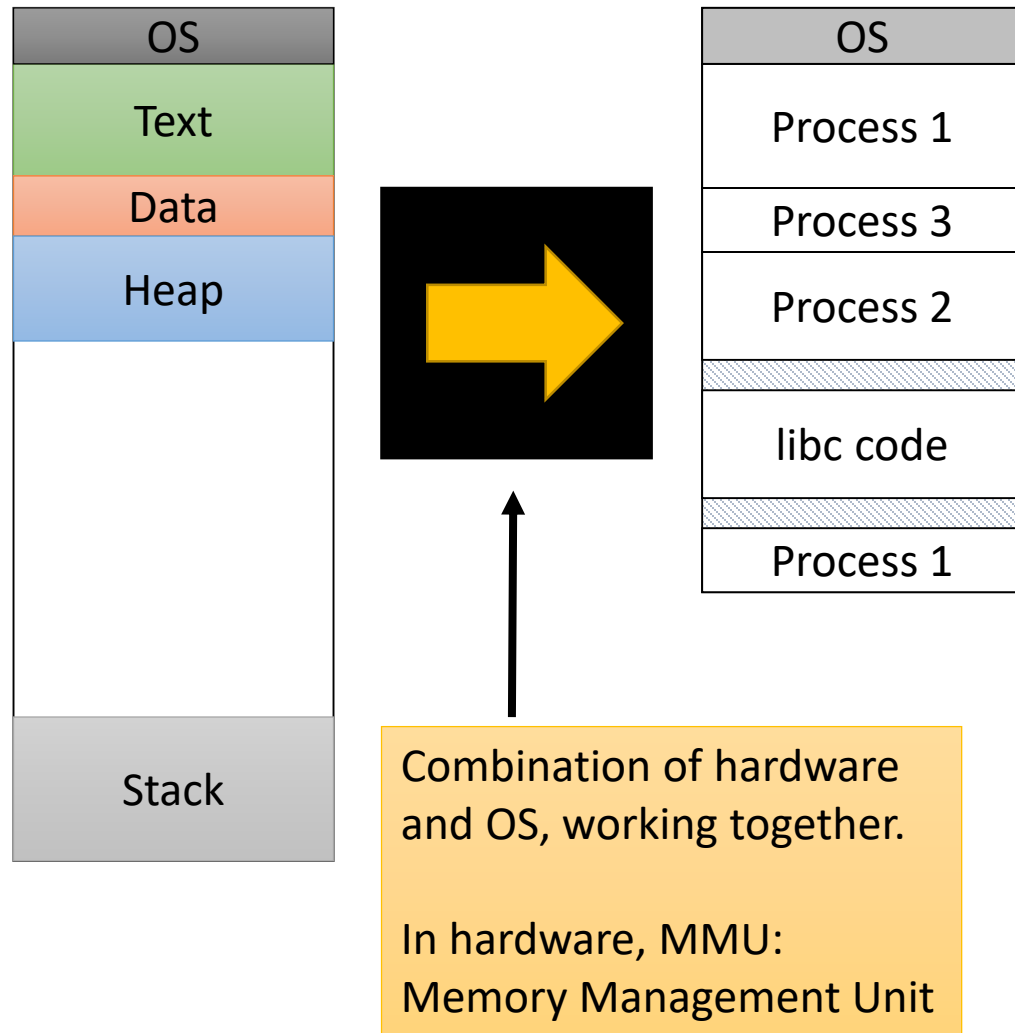
Swarthmore College

March 8, 2018

Today's Goals

- Describe the mechanisms behind address translation.
- Analyze the performance of address translation alternatives.
- ~~Explore page replacement policies for disk swapping.~~

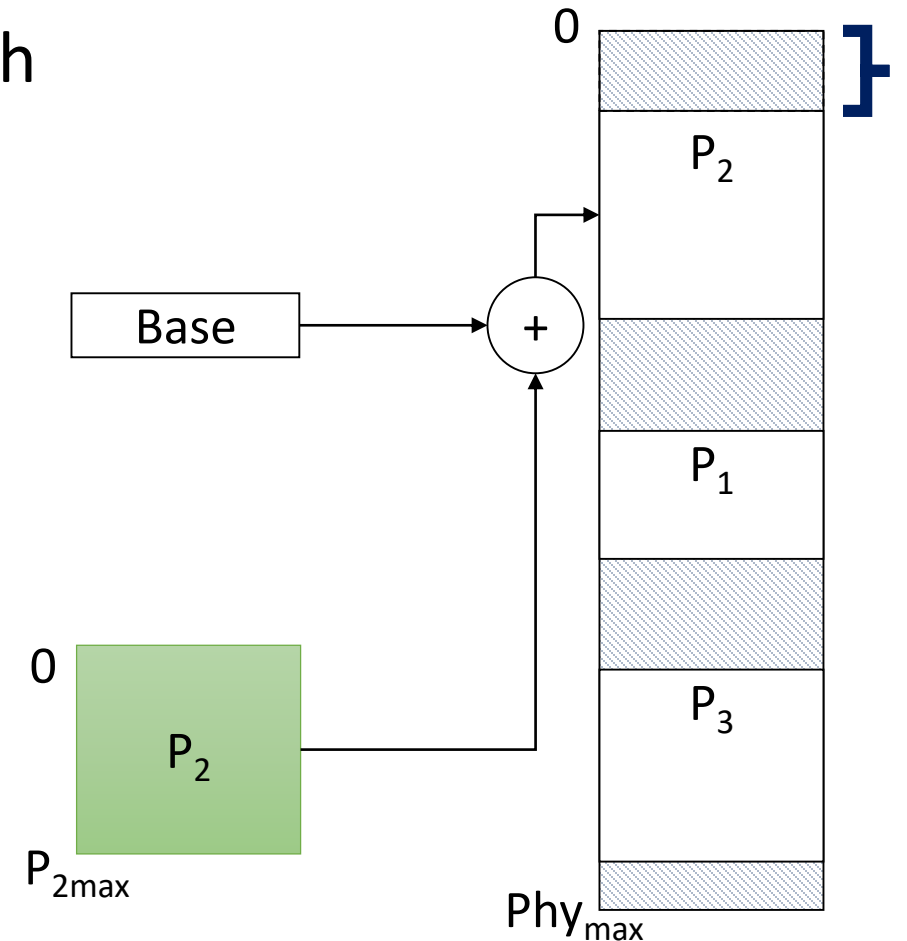
Address Translation: Wish List



- Map virtual addresses to physical addresses.
- Allow multiple processes to be in memory at once, but isolate them from each other.
- Determine which subset of data to keep in memory / move to disk.
- Allow the same physical memory to be mapped in multiple process VASes.
- Make it easier to perform placement in a way that reduces fragmentation.
- Map addresses quickly with a little HW help.

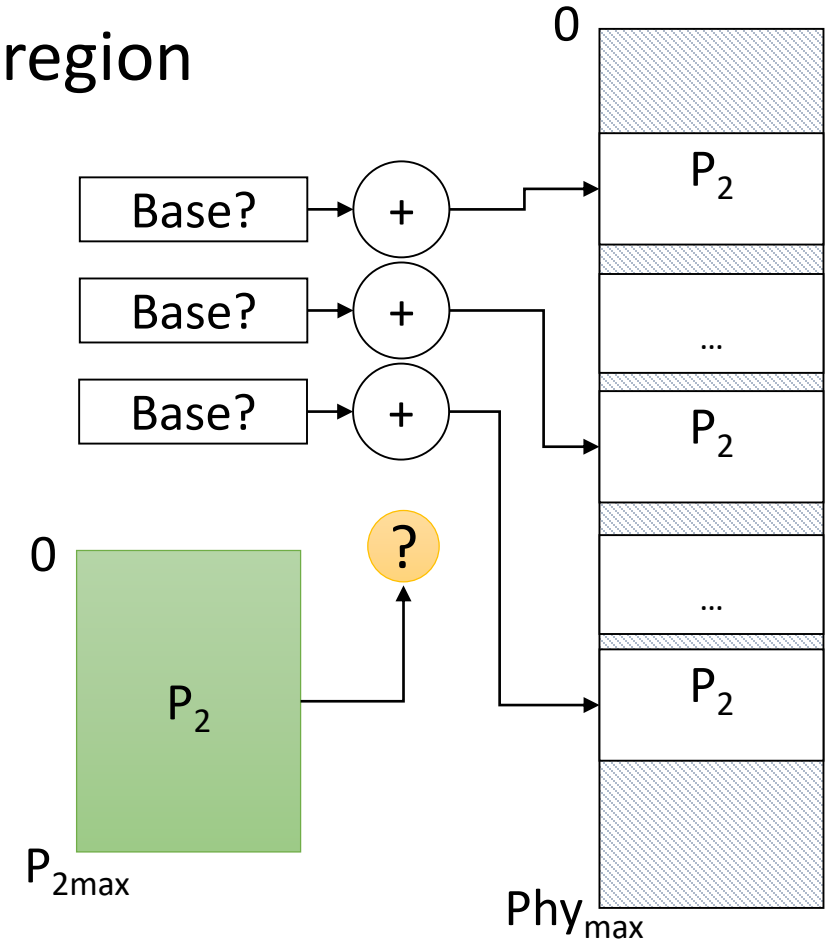
Simple (Unrealistic) Translation Example

- Process P_2 's virtual addresses don't align with physical memory's addresses.
- Determine offset from physical address 0 to start of P_2 , store in *base*.



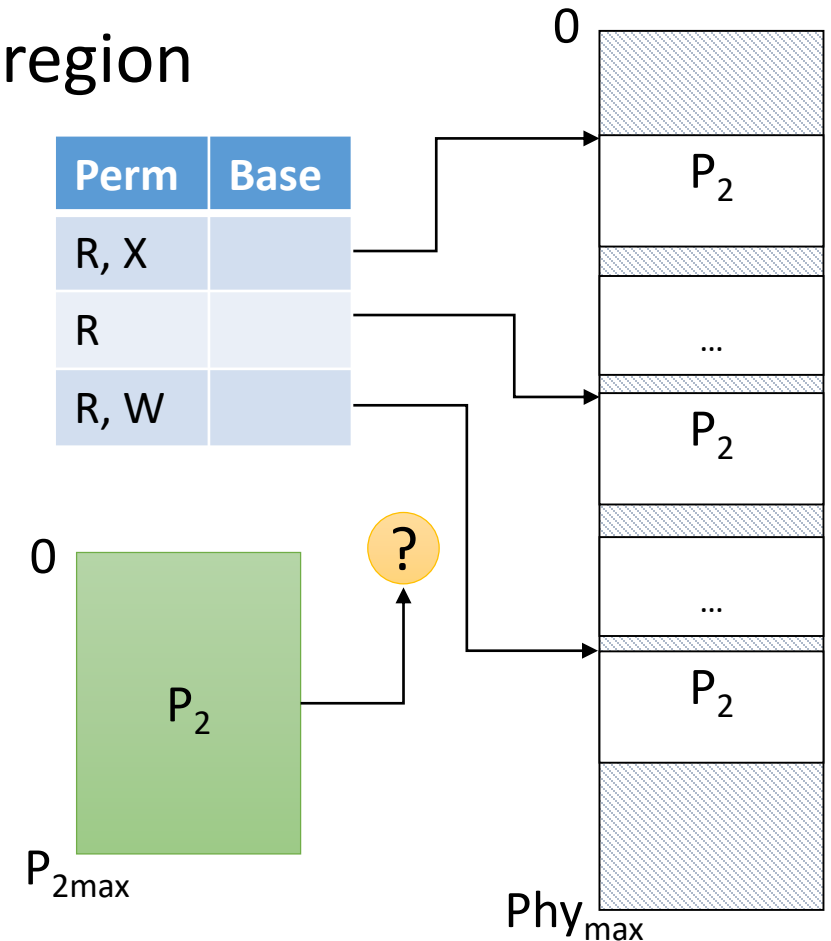
Generalizing

- Problem: process may not fit in one contiguous region




Generalizing

- Problem: process may not fit in one contiguous region
- Solution: keep a table (one for each process)
 - Keep details for each region in a row
 - Store additional metadata (ex. permissions)
- Interesting questions:
 - How many regions should there be (and what size)?
 - How to determine which row we should use?




Defining Regions - Two Approaches

- Segmentation:
 - Partition address space and memory into segments
 - Segments have varying sizes




- Paging:
 - Partition address space and memory into pages
 - Pages are a constant, fixed size



Fragmentation

Internal


- Process asks for memory, doesn't use it all.
- Possible reasons:
 - Process was wrong about needs
 - OS gave it more than it asked for
- *internal*: within an allocation



Memory allocated to process

External

- Over time, we end up with these small gaps that become more difficult to use (eventually, wasted).
- *external*: unused memory between allocations




Which scheme is better for reducing internal and external fragmentation.

- A. Segmentation is better than paging for both forms of fragmentation.
- B. Segmentation is better for *internal* fragmentation, and paging is better for *external* fragmentation.
- C. Paging is better for *internal* fragmentation, and segmentation is better for *external* fragmentation.
- D. Paging is better than segmentation for both forms of fragmentation.

Which would you use? Why? Pros/Cons?


A. Segmentation:

- Partition address space and memory into segments
- Segments have varying sizes



B. Paging:

- Partition address space and memory into pages
- Pages are a constant, fixed size




C. Something else (what?)

Segmentation vs. Paging

- A segment is good *logical* unit of information
 - Can be sized to fit any contents
 - Easy to share large regions (e.g., code, data)
 - Protection requirements correspond to logical data segment
- A page is good *physical* unit of information
 - Simple physical memory placement
 - No external fragmentation
 - Constant sizes make it easier for hardware to help

Generalizing

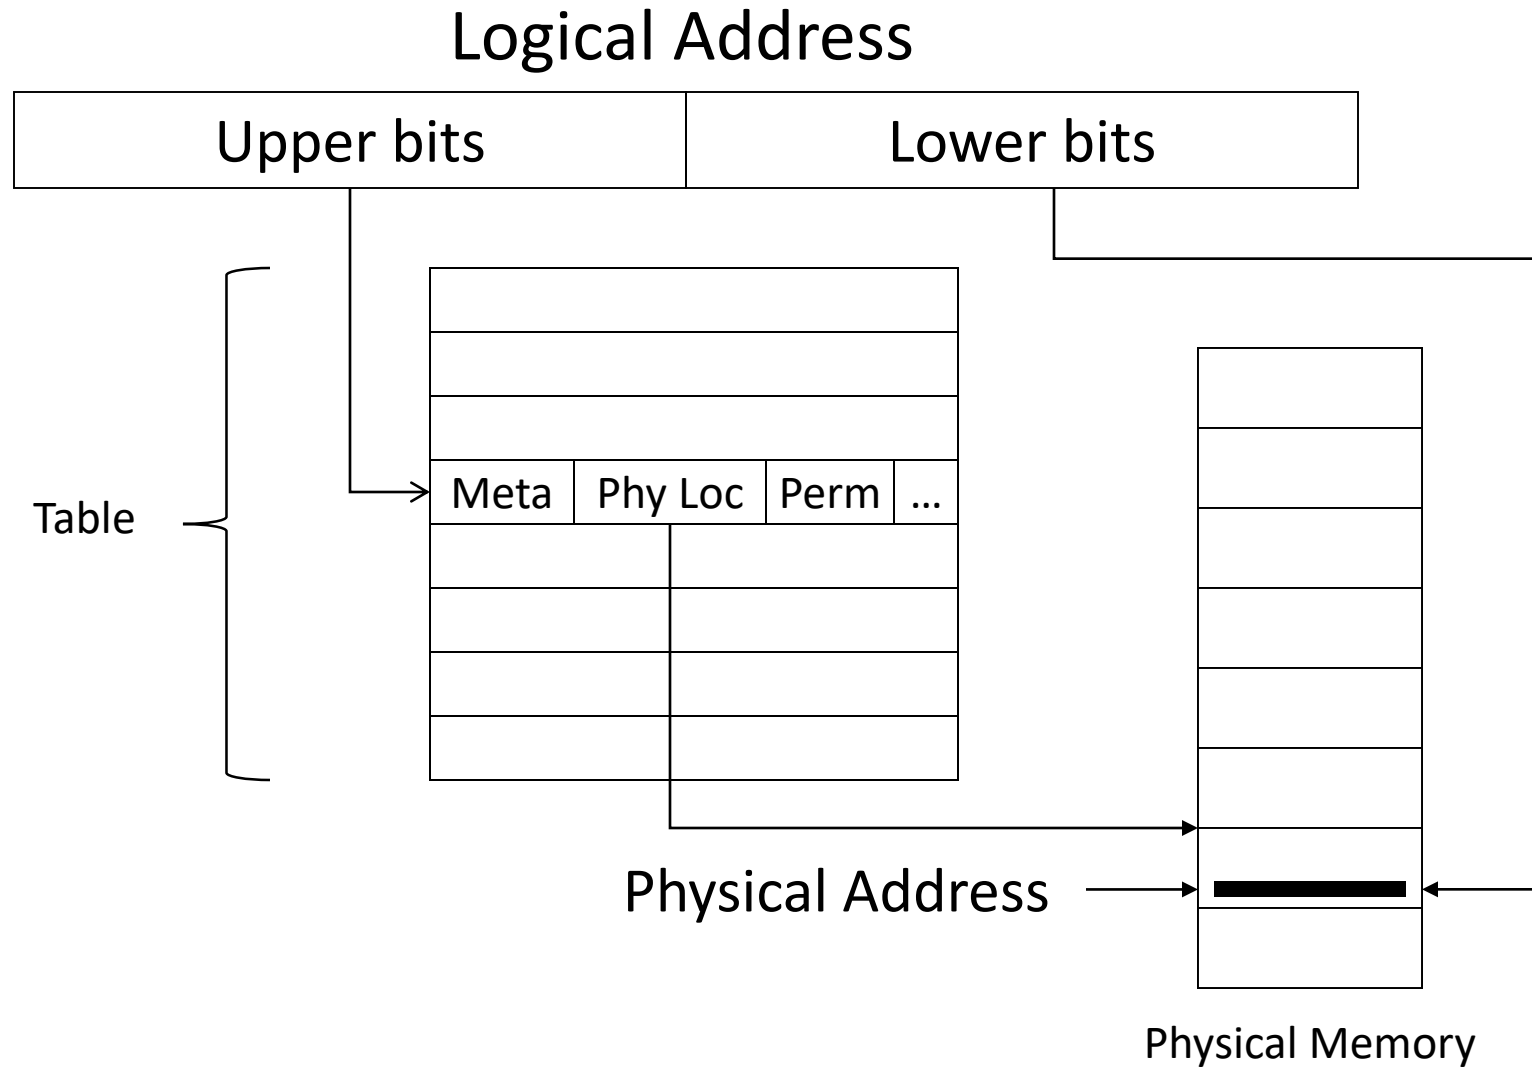
- Problem: process may not fit in one contiguous region
- Solution: keep a table (one for each process)
 - Keep details for each region in a row
 - Store additional metadata (ex. permissions)
- Interesting questions:
 - How many regions should there be (and what size)?
 - **How to determine which row we should use?**



For **both** segmentation and paging...

- Each process gets a table to track memory address translations.
- When a process attempts to read/write to memory:
 - use high order bits of virtual address to determine which row to look at in the table
 - use low order bits of virtual address to determine an offset within the physical region

Address Translation



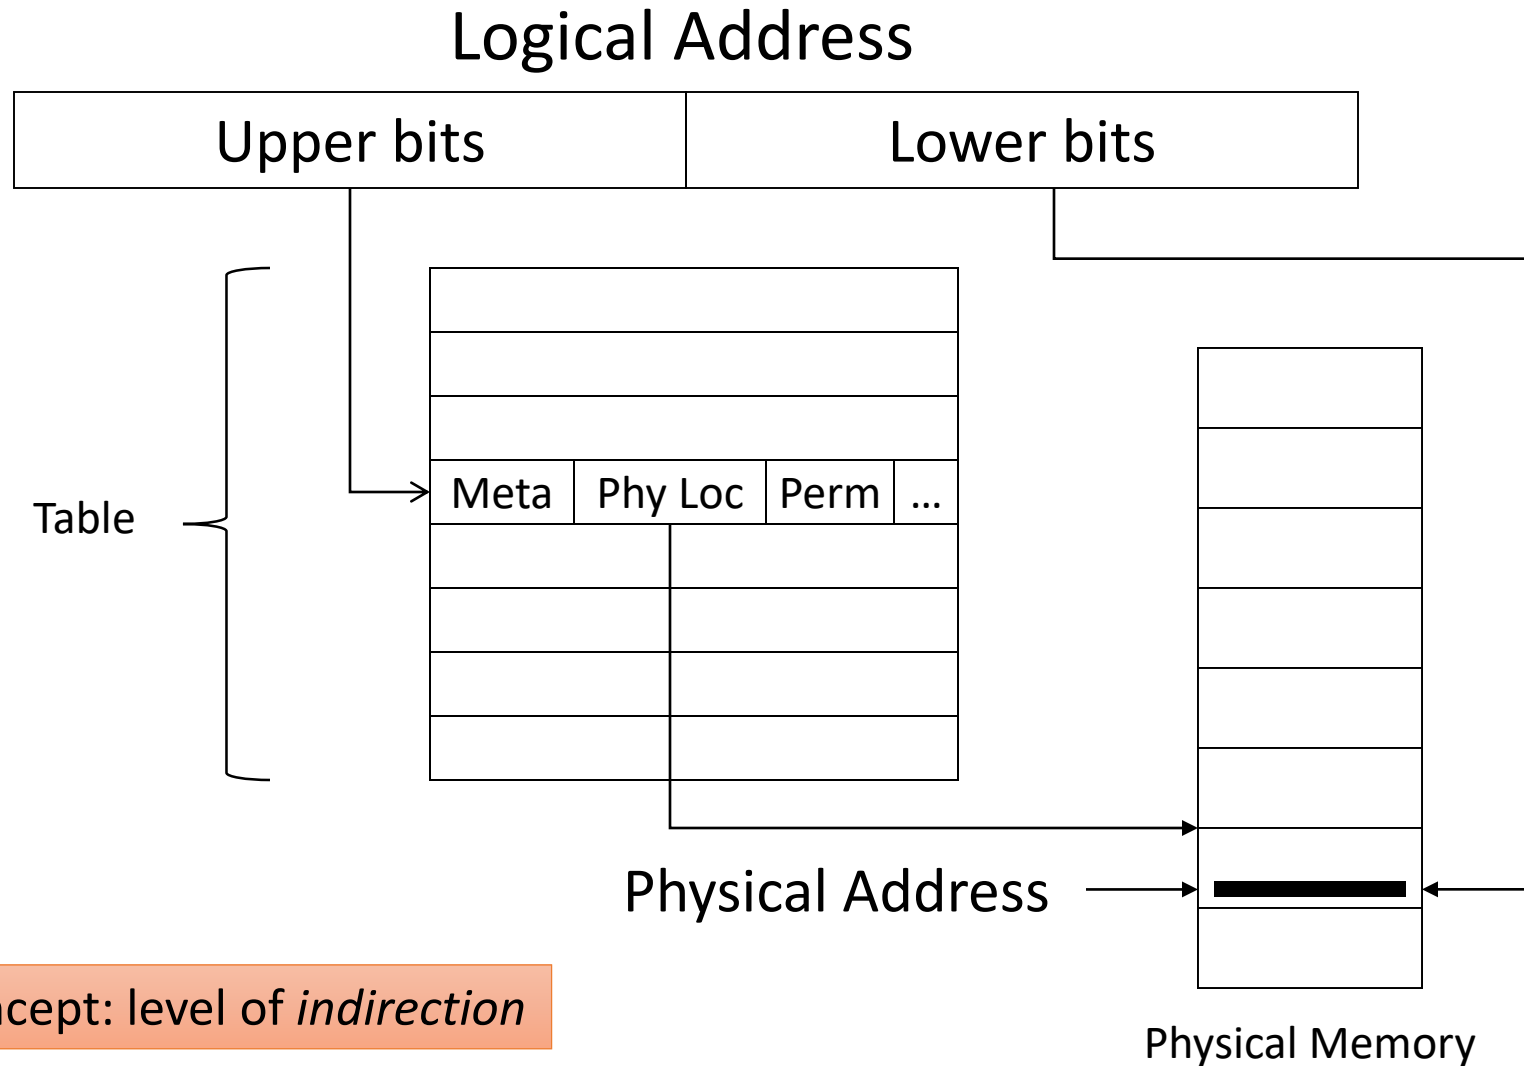
Performance Implications

Without VM:

Go directly to address in memory.


With VM:

Do a lookup in memory to determine which address to use.




Defining Regions - Two Approaches

- Segmentation:
 - Partition address space and memory into segments
 - Segments have varying sizes

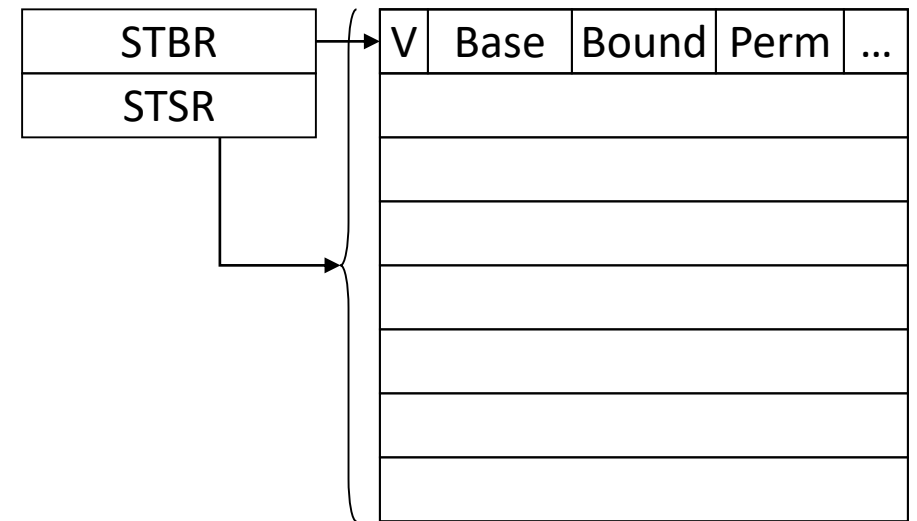


- Paging:
 - Partition address space and memory into pages
 - Pages are a constant, fixed size

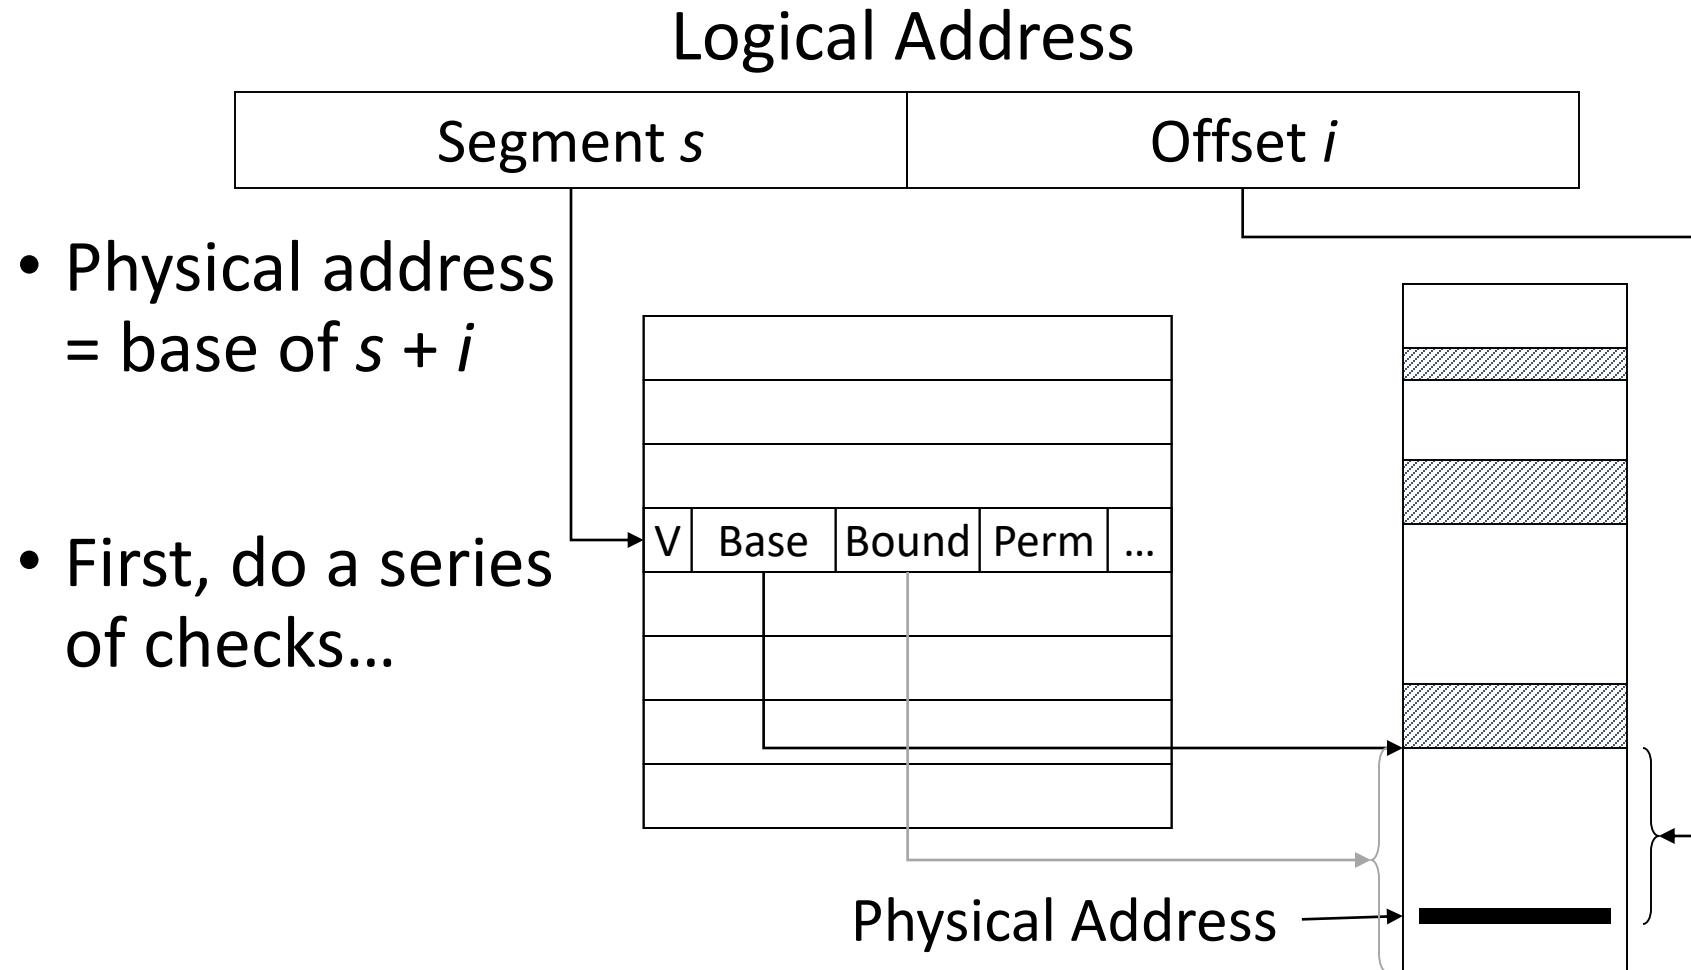


Segment Table

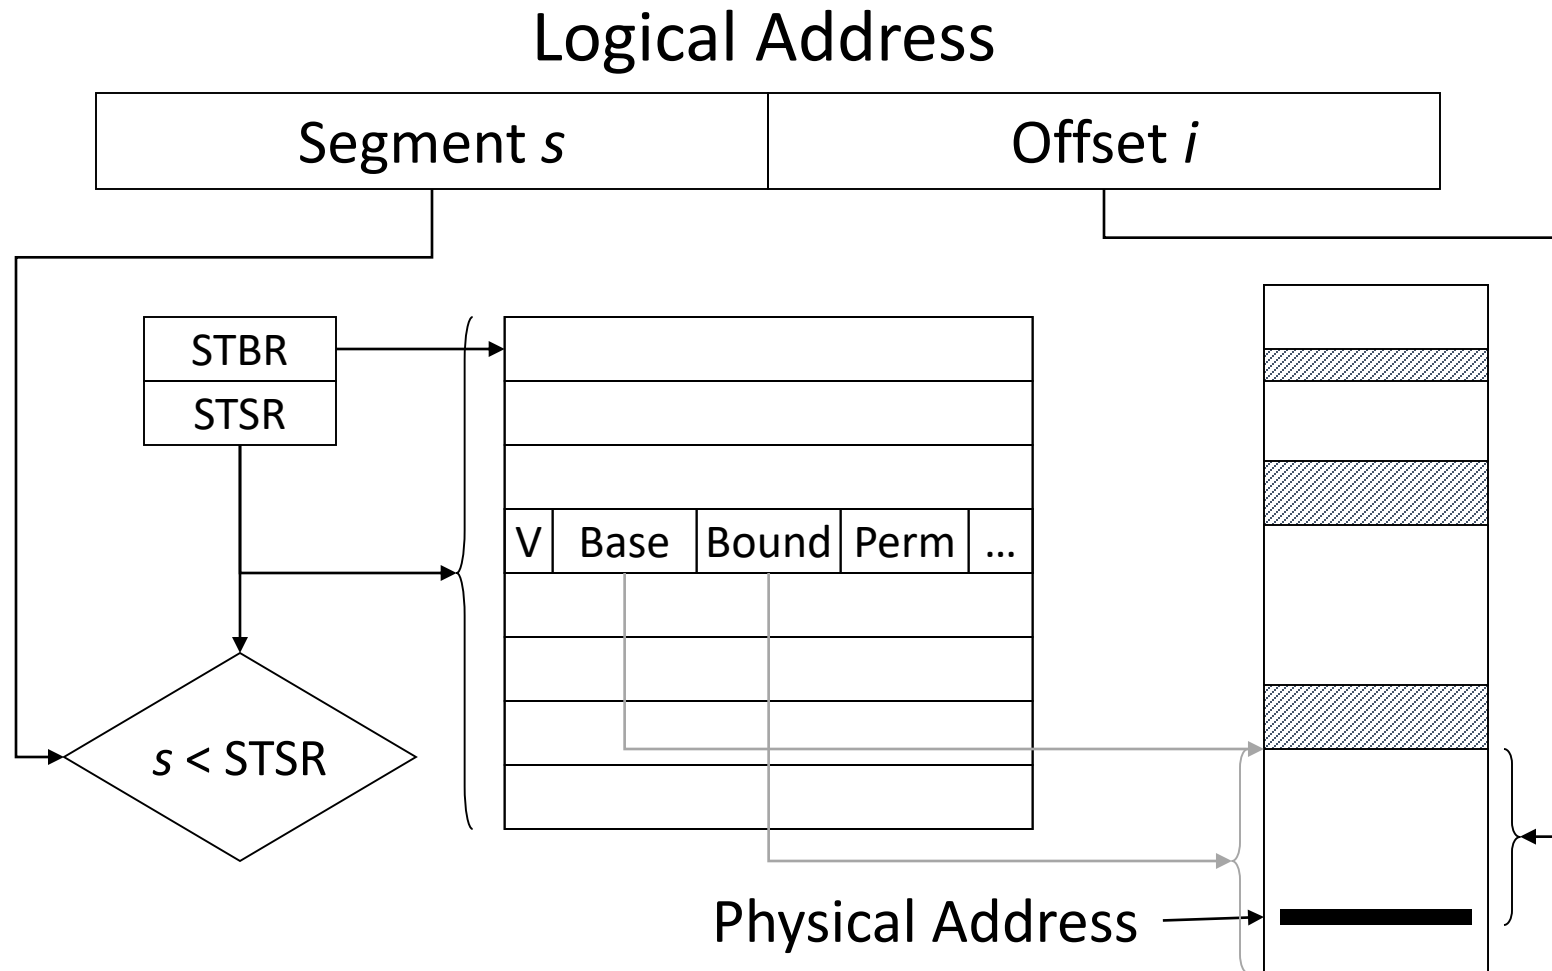
- One table per process
- Where is the *table* located in memory?
 - Segment table base register (STBR)
 - Segment table size register (STSR)
- Table entry elements
 - V: valid bit (does it contain a mapping?)
 - Base: segment location in physical memory
 - Bound: segment size in physical memory
 - Permissions



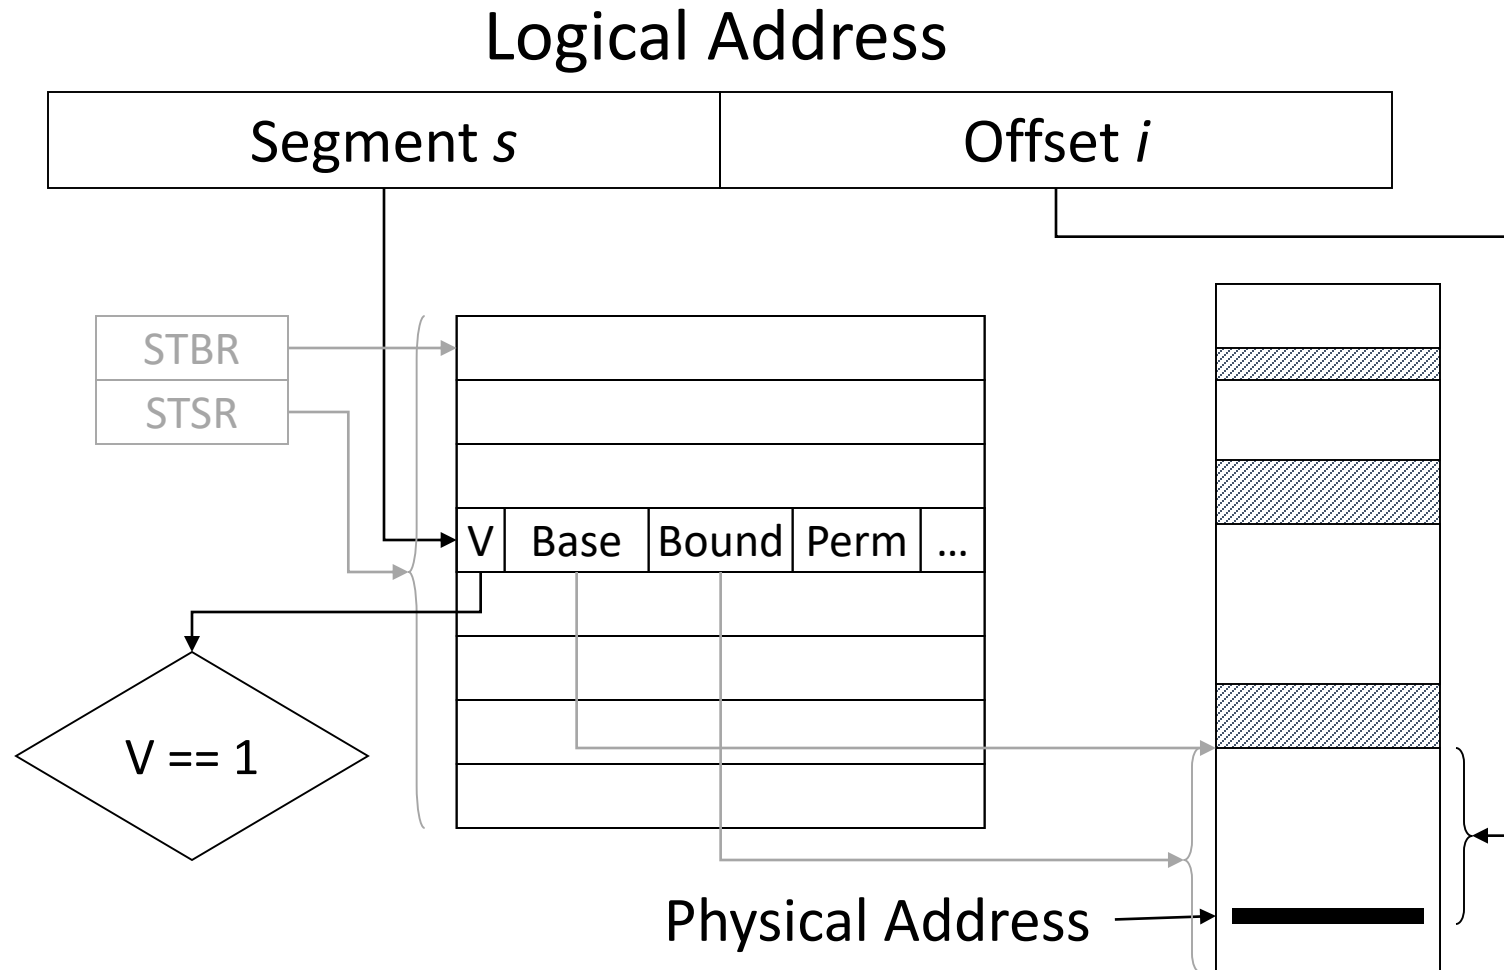
Address Translation



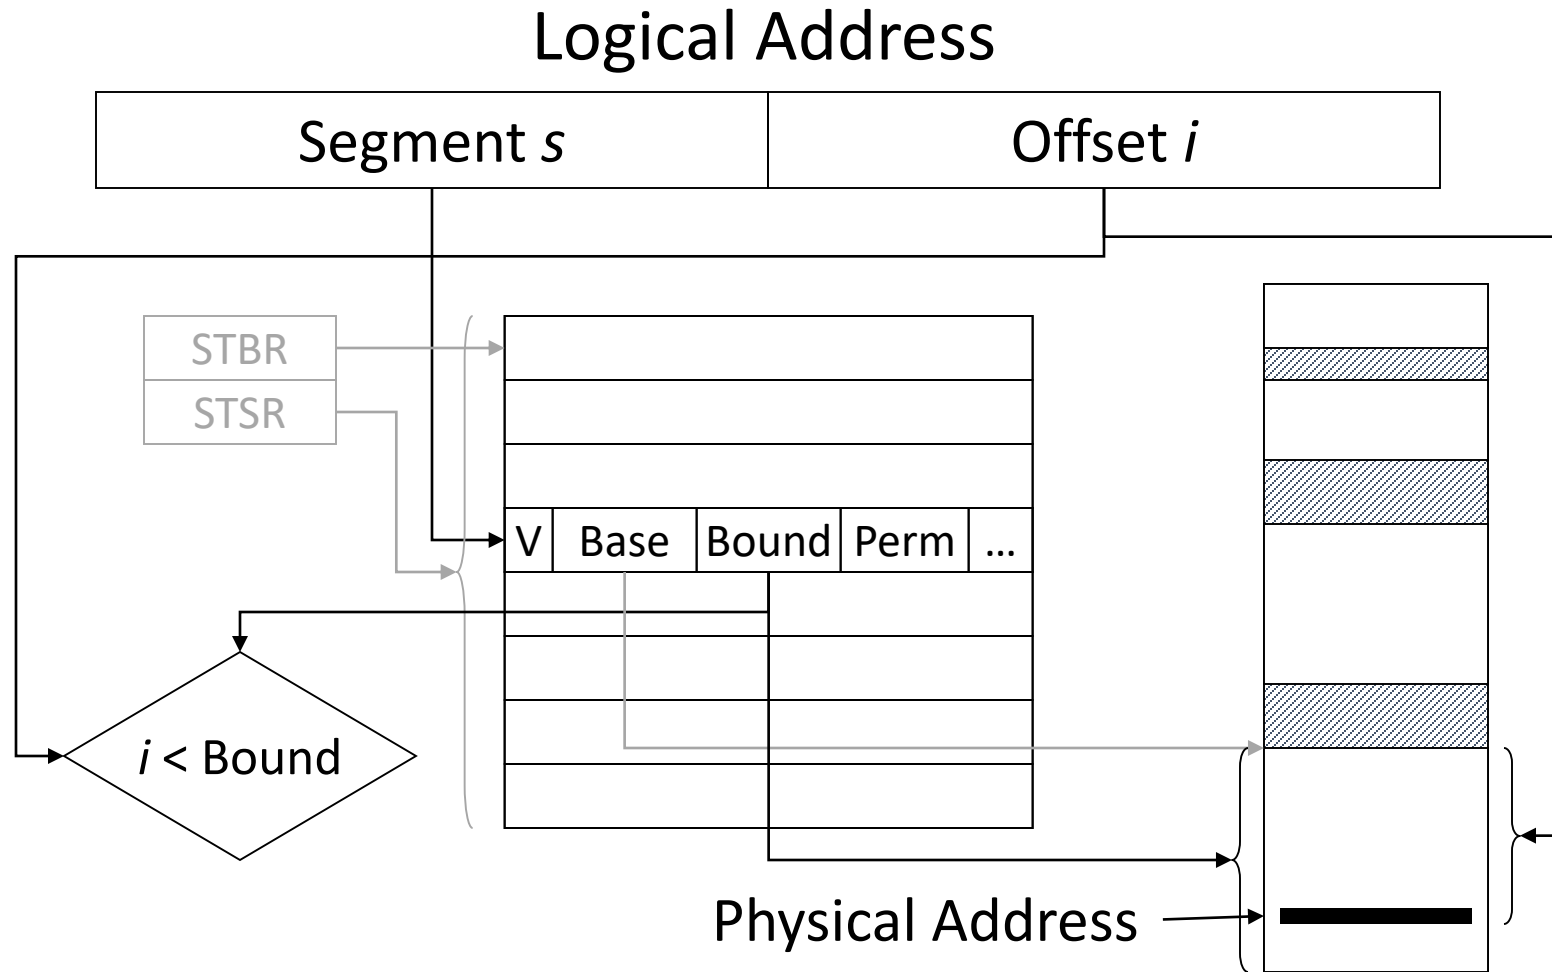
Check if Segment s is within Range



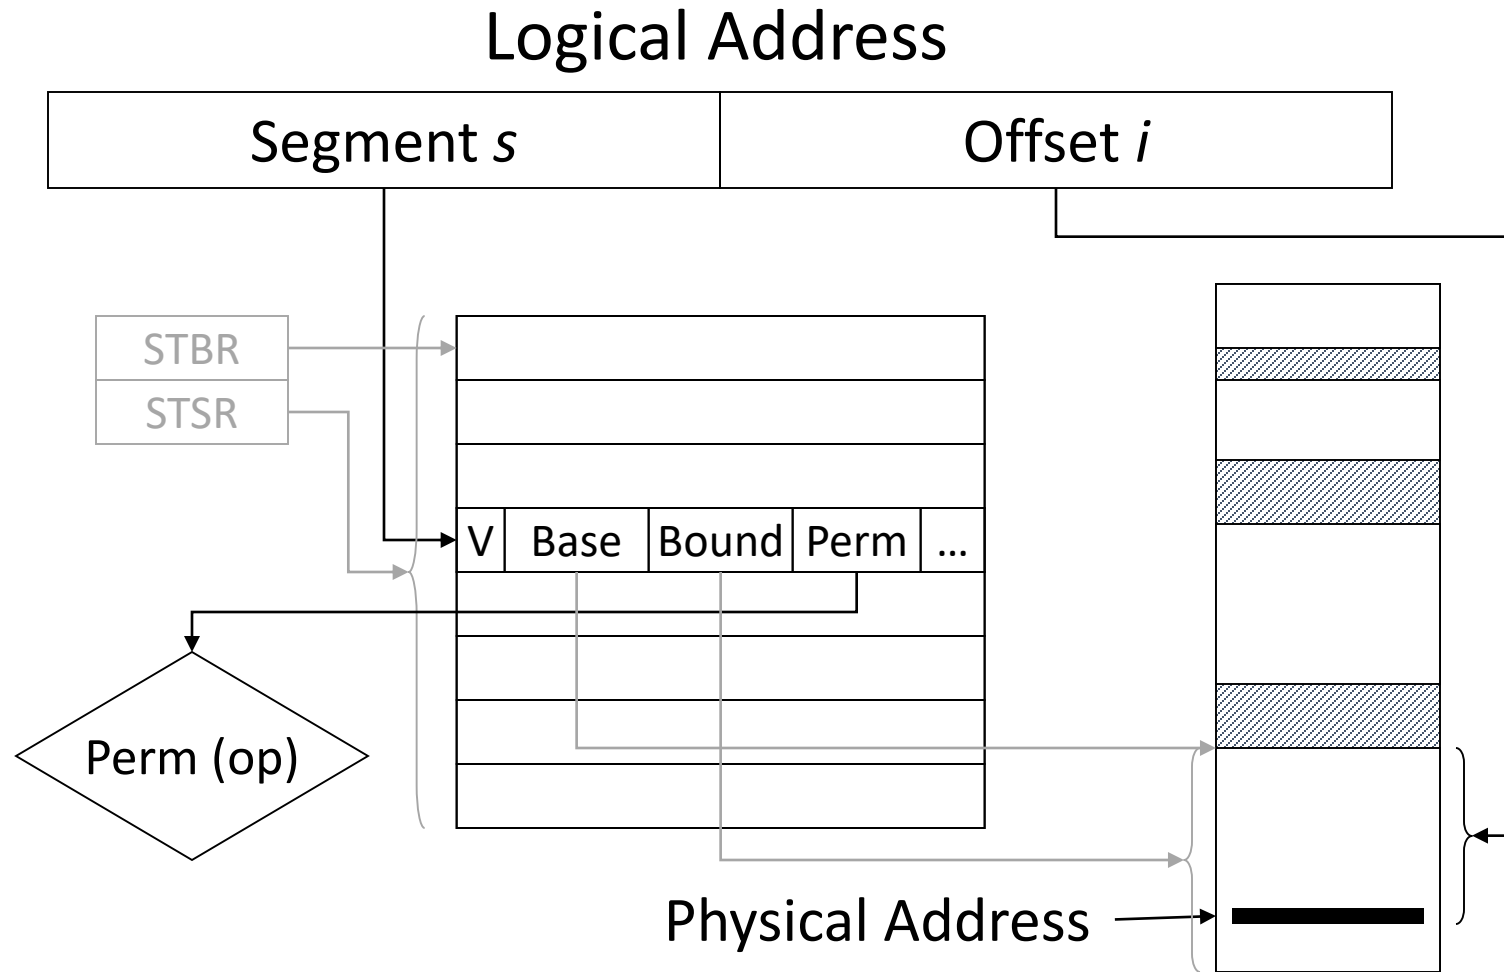
Check if Segment Entry s is Valid



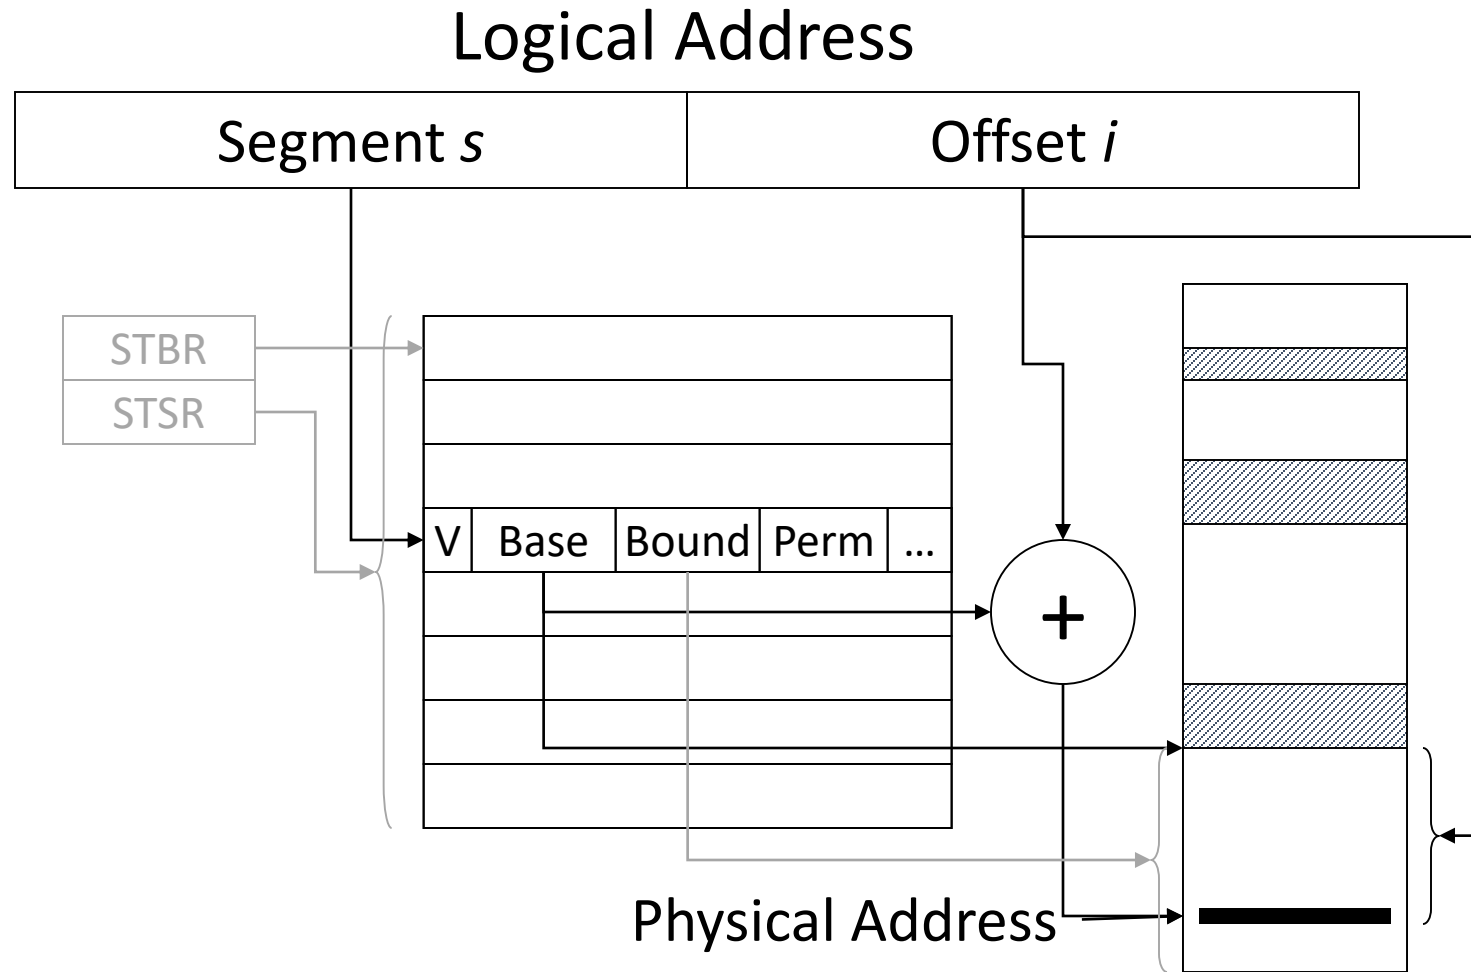
Check if Offset i is within Bounds



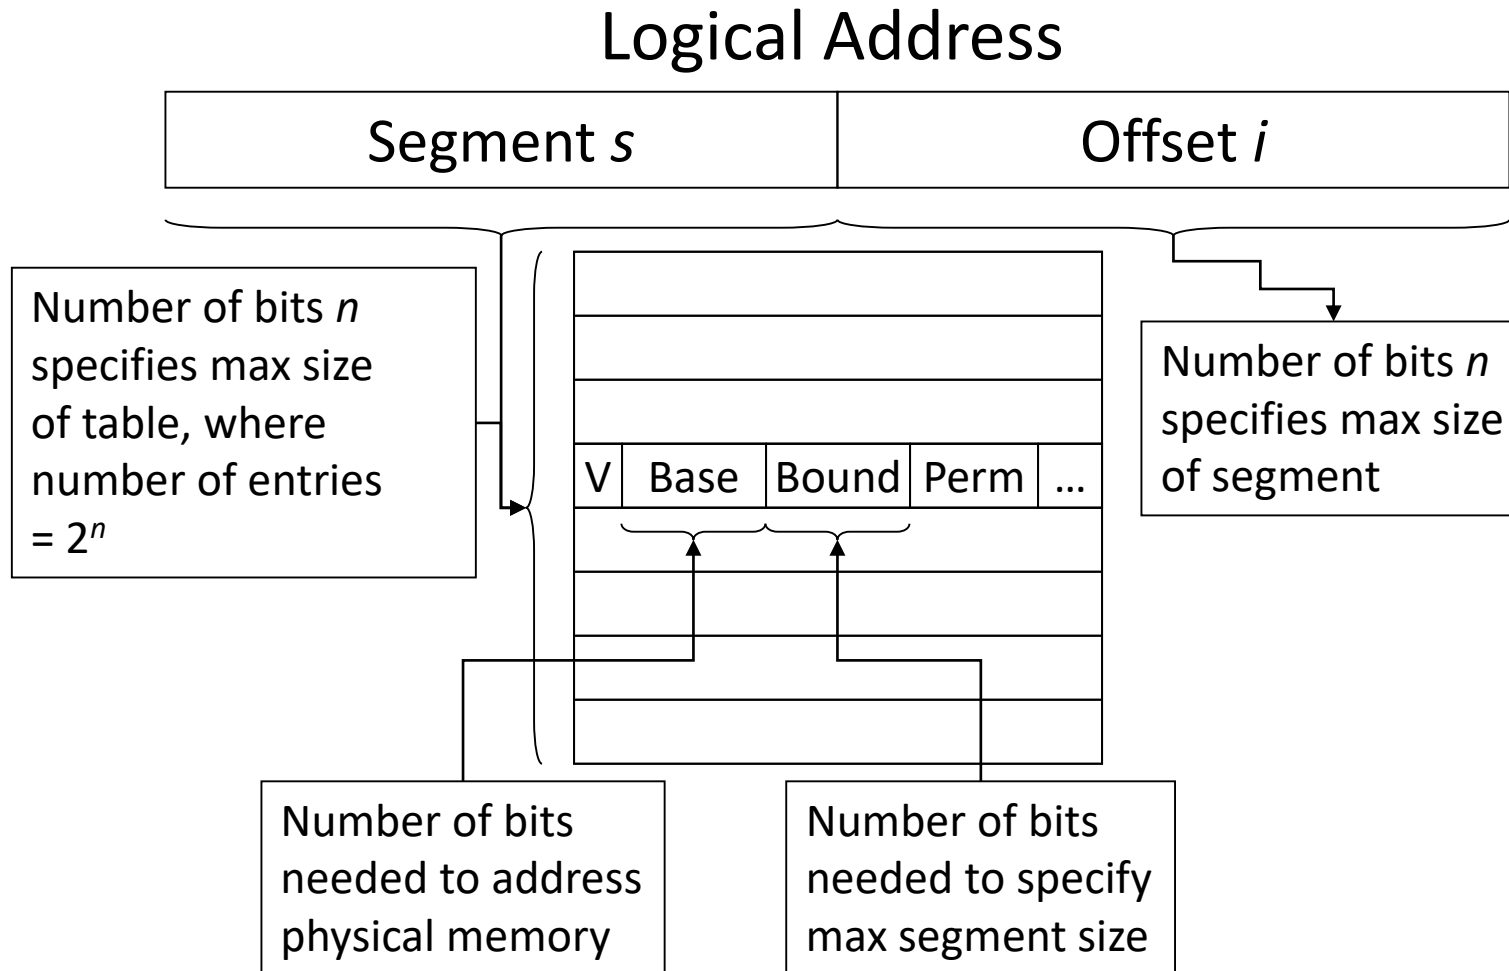
Check if Operation is Permitted



Translate Address



Sizing the Segment Table



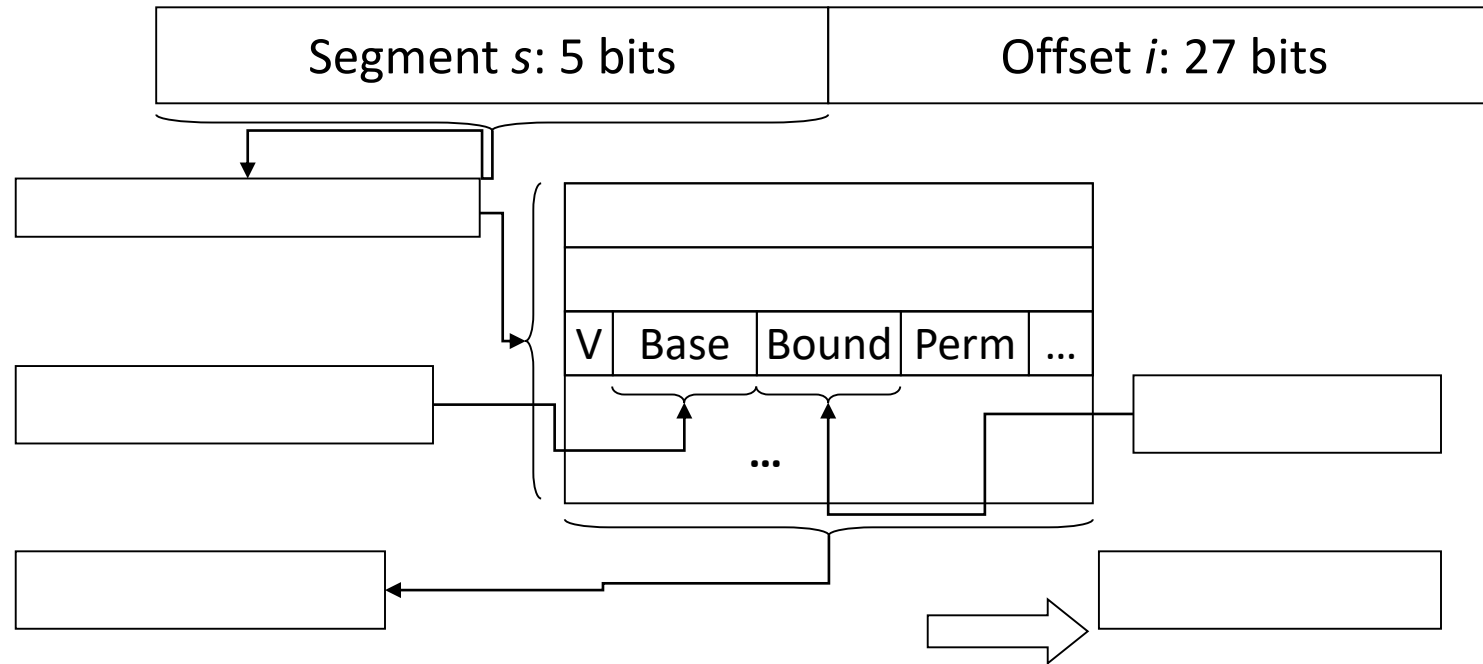
Helpful reminder:

$2^{10} \Rightarrow$ Kilobyte

$2^{20} \Rightarrow$ Megabyte

$2^{30} \Rightarrow$ Gigabyte

Example of Sizing the Segment Table



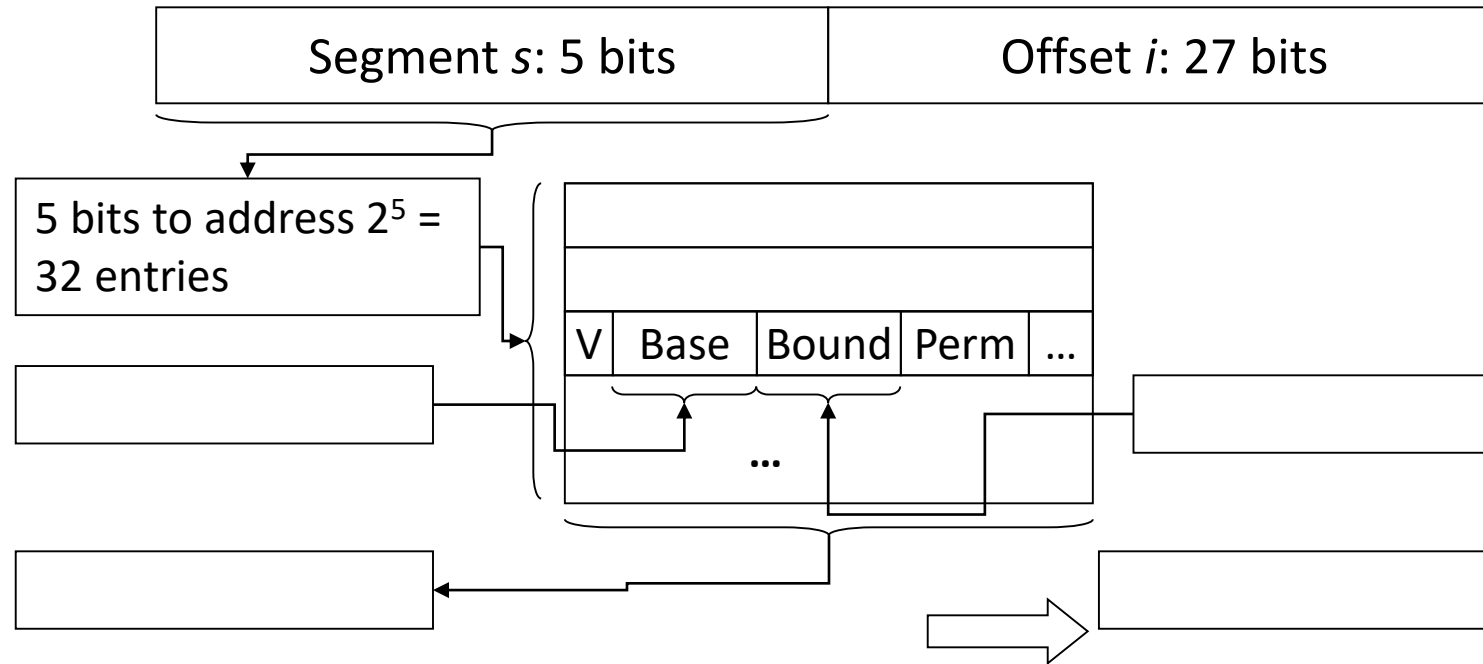
- Given 32-bit virtual address space, 1 GB physical memory (max)
 - 5 bit segment number, 27 bit offset

5 bit segment address, 32 bit logical address, 1 GB Physical memory.

How many entries (rows) will we have in our segment table?

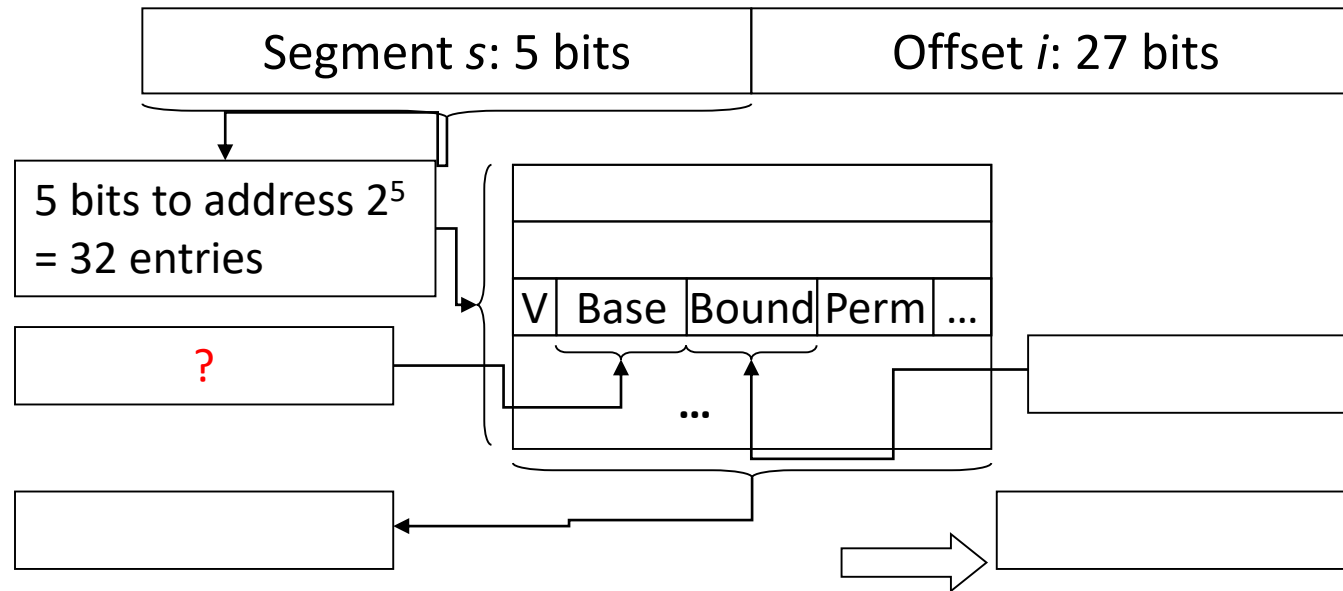
- A. 32: The logical address size is 32 bits
- B. 32: The segment address is five bits
- C. 30: We need to address 1 GB of physical memory
- D. 27: We need to address up to the maximum offset

Example of Sizing the Segment Table



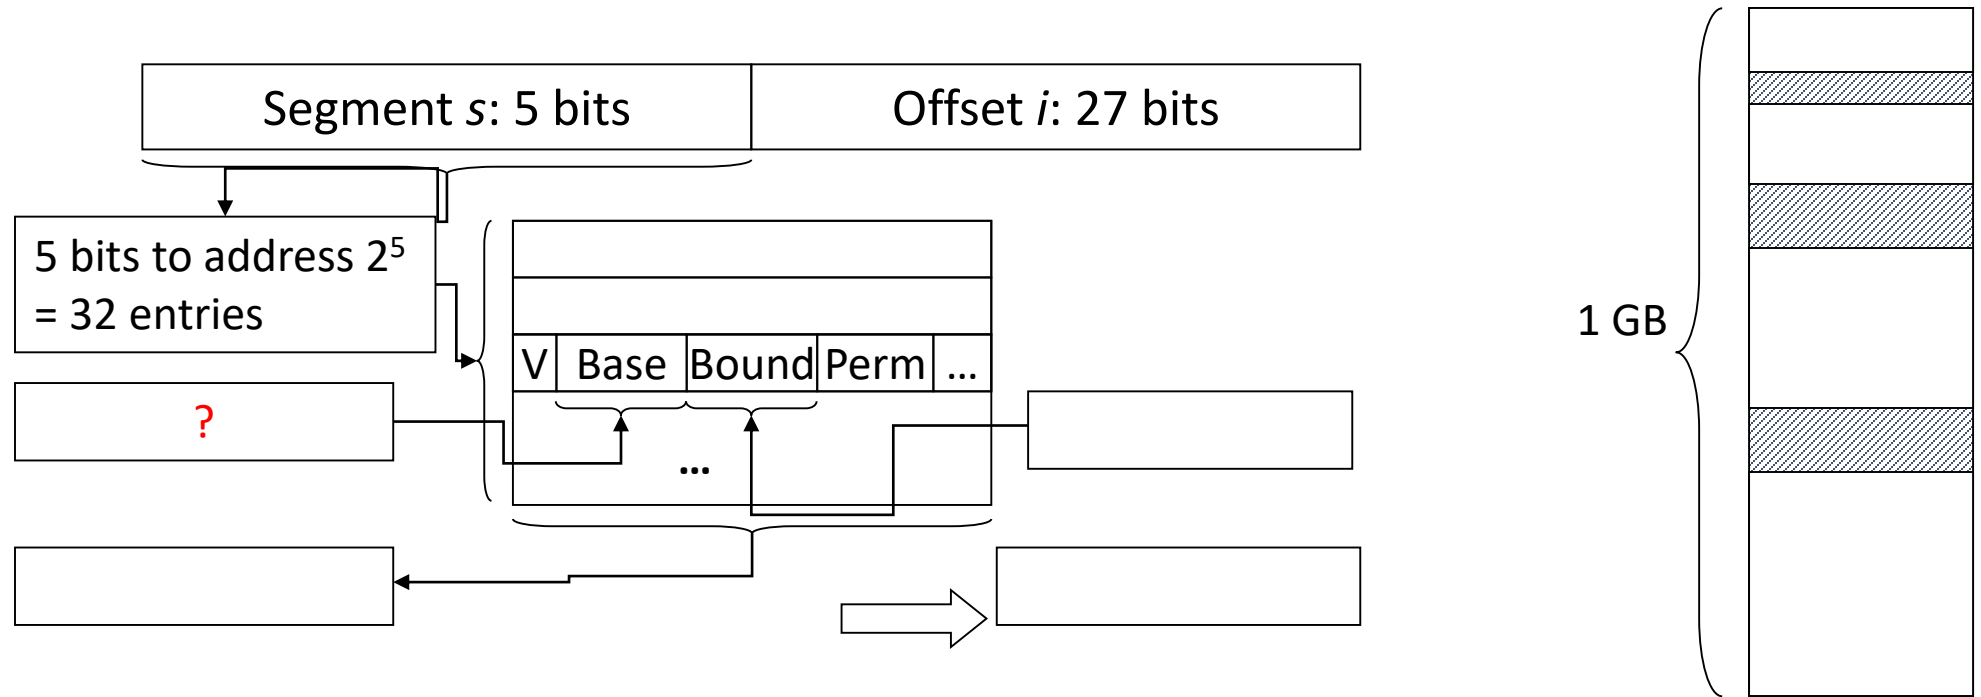
- Given 32-bit virtual address space, 1 GB physical memory (max)
 - 5 bit segment number, 27 bit offset

How many bits do we need for the base?



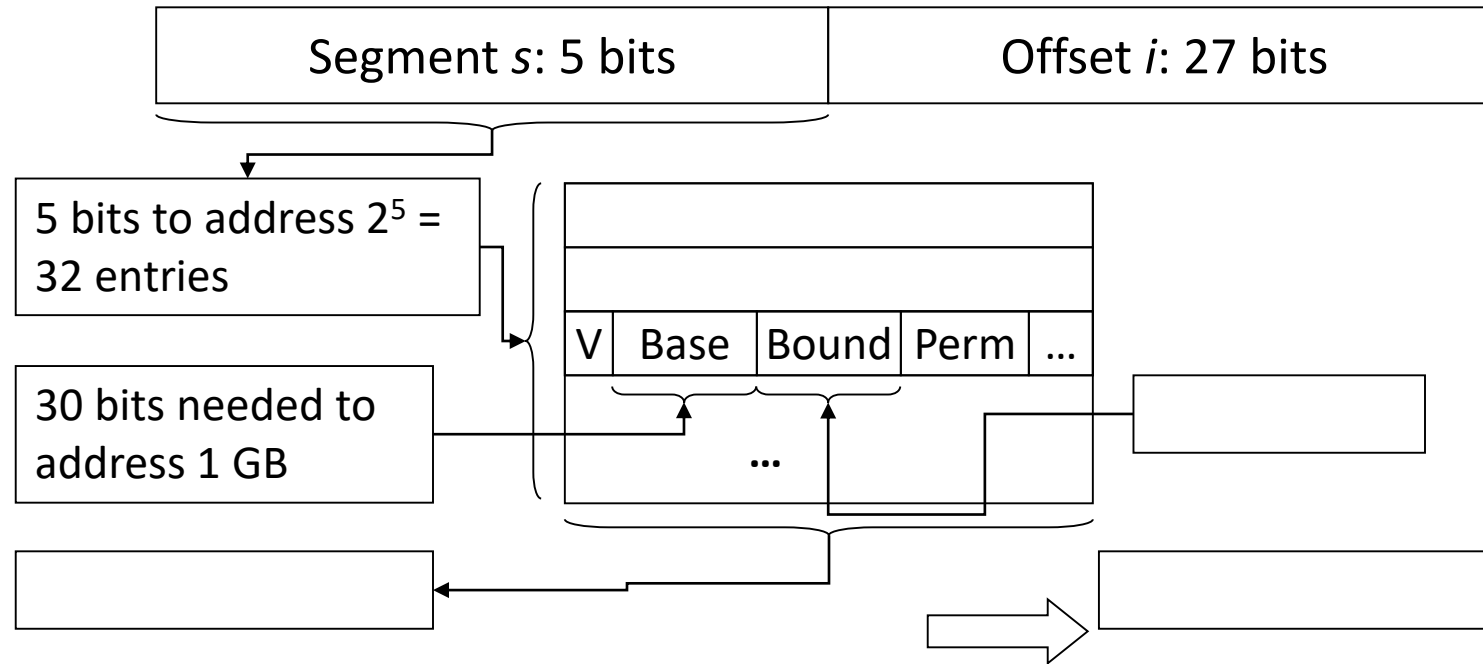
- A. 30 bits, to address 2^{30} bytes (1 GB) of physical memory
- B. 5 bits, because that's how many the segment bits we have
- C. 27 bits, because that's how many offset bits we have

How many bits do we need for the base?



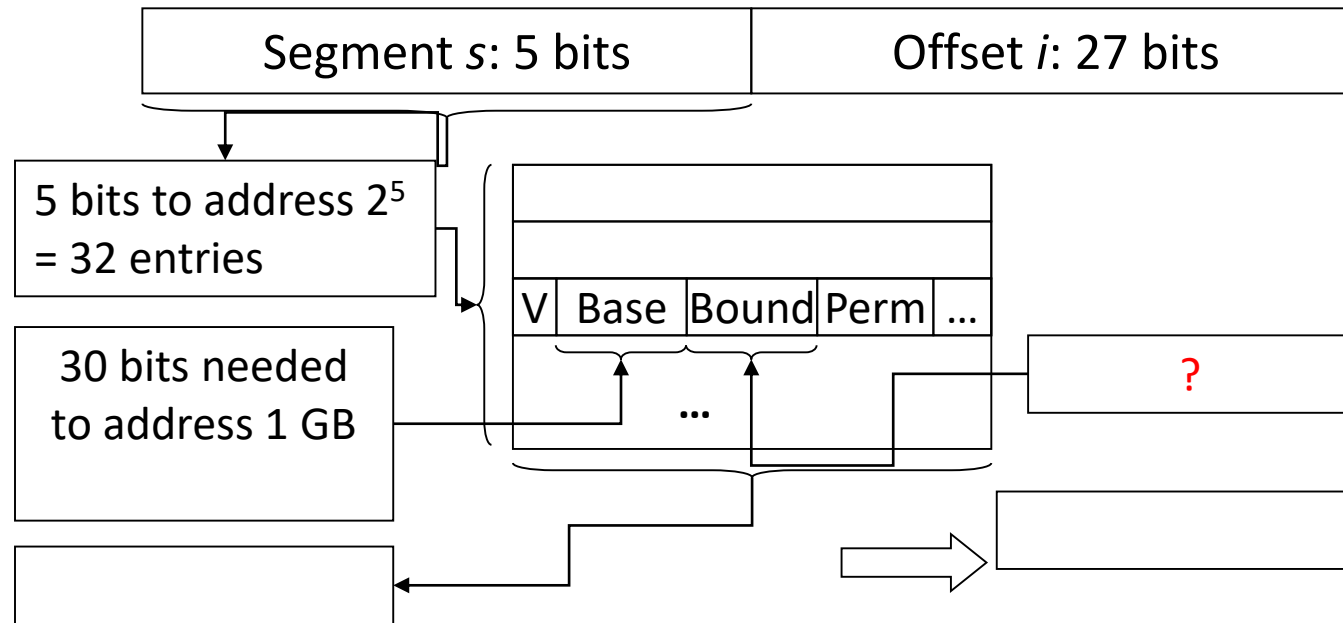
- A. 30 bits, to address 2^{30} bytes (1 GB) of physical memory
- B. 5 bits, because that's how many the segment bits we have
- C. 27 bits, because that's how many offset bits we have

Example of Sizing the Segment Table



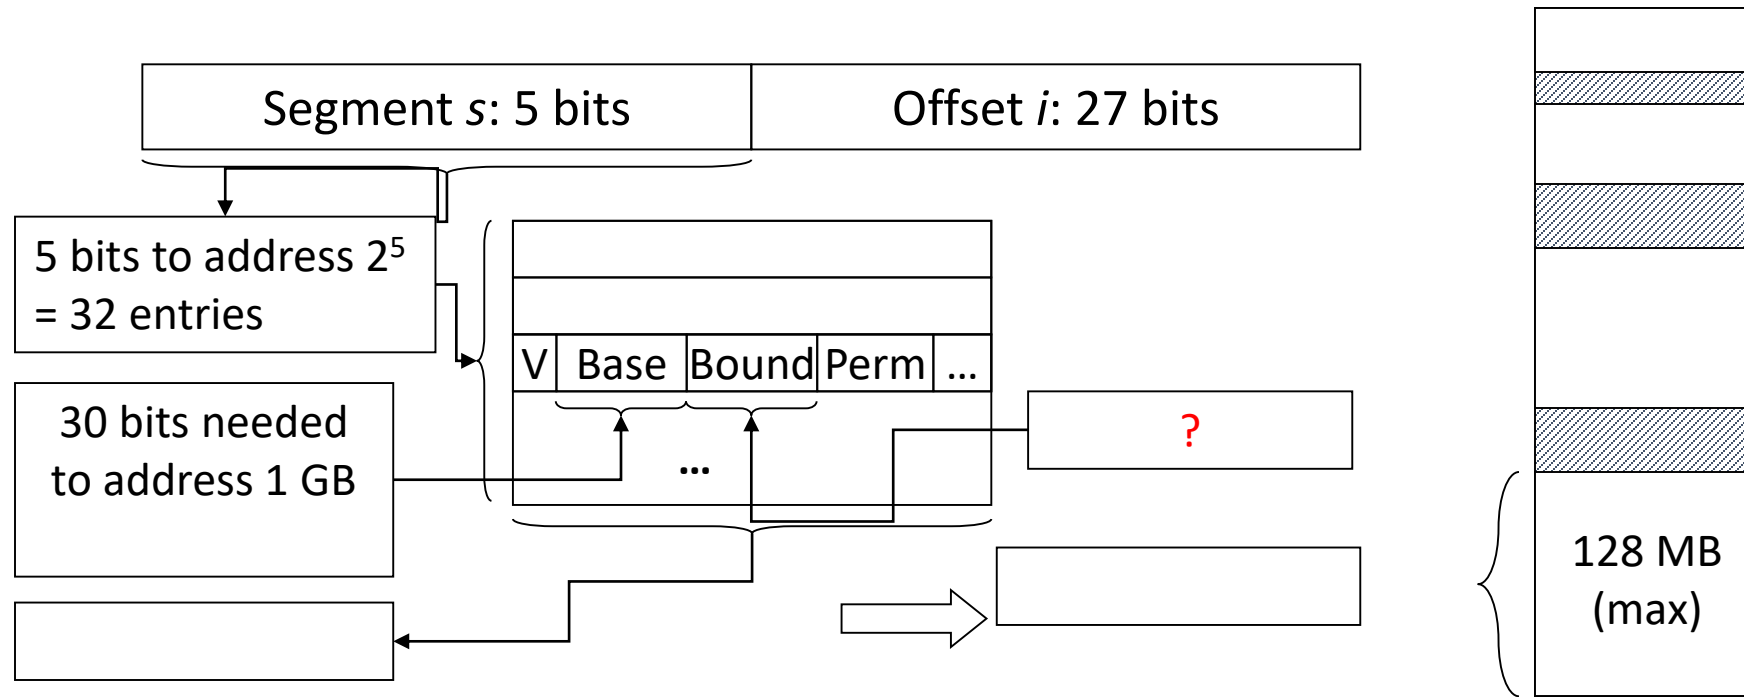
- Given 32-bit virtual address space, 1 GB physical memory (max)
 - 5 bit segment number, 27 bit offset

How many bits do we need for the bound?



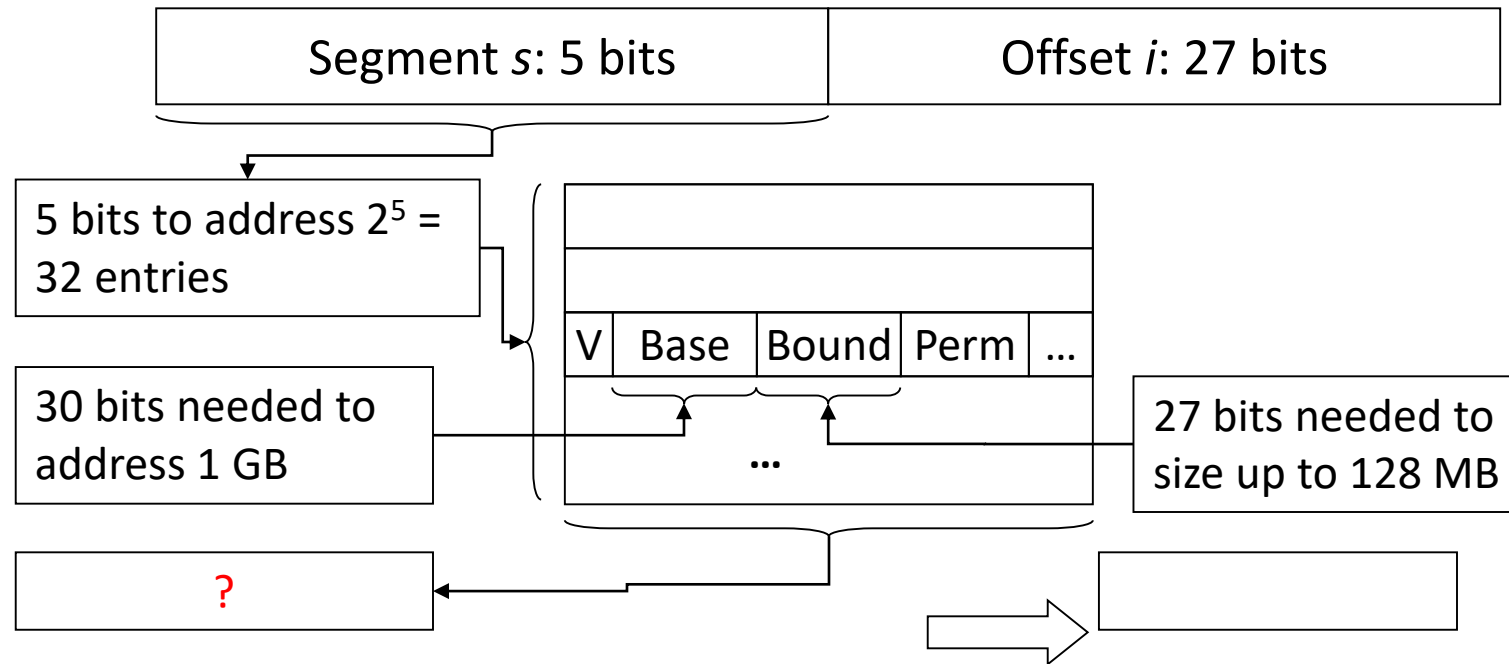
- A. 5 bits: the size of the segment portion of the virtual address
- B. 27 bits: the size of the offset portion of the virtual address
- C. 32 bits: the size of the virtual address

How many bits do we need for the bound?



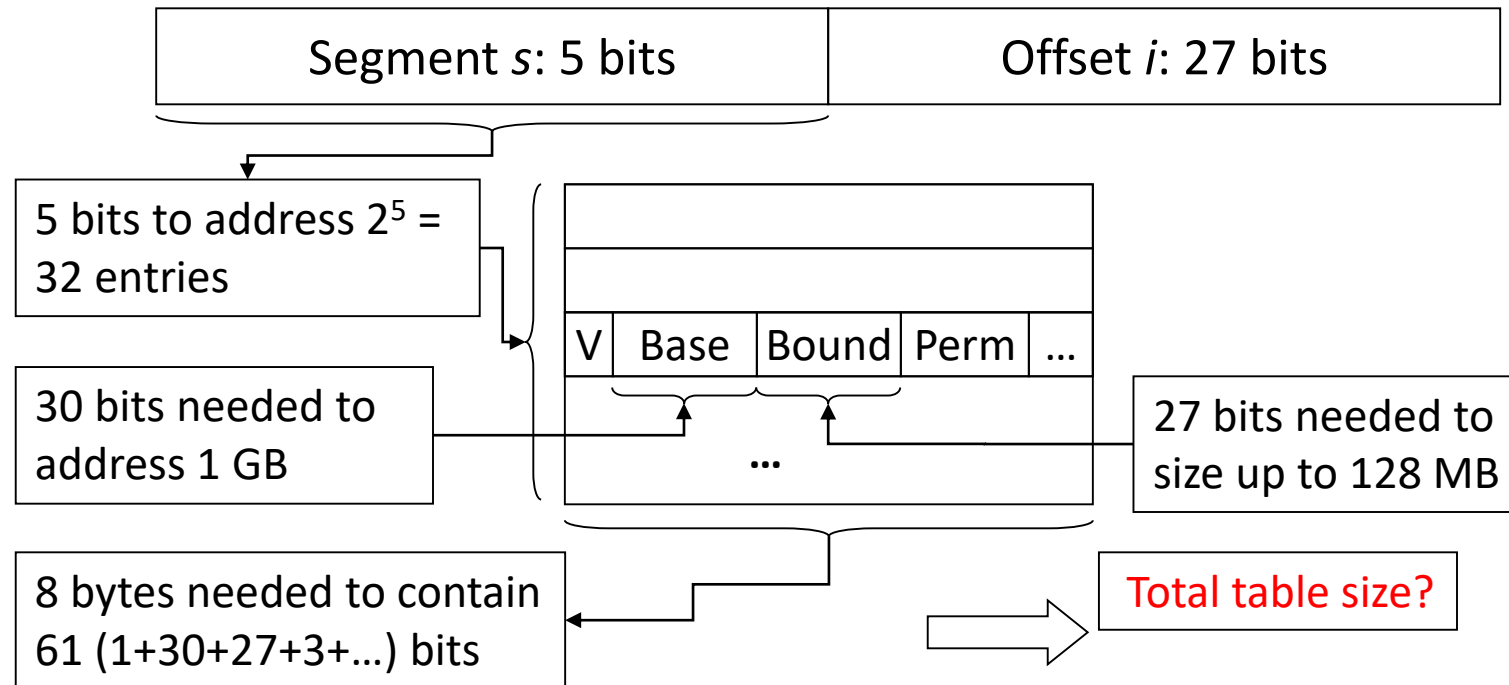
- A. 5 bits: the size of the segment portion of the virtual address
- B. 27 bits: the size of the offset portion of the virtual address
- C. 32 bits: the size of the virtual address

Example of Sizing the Segment Table



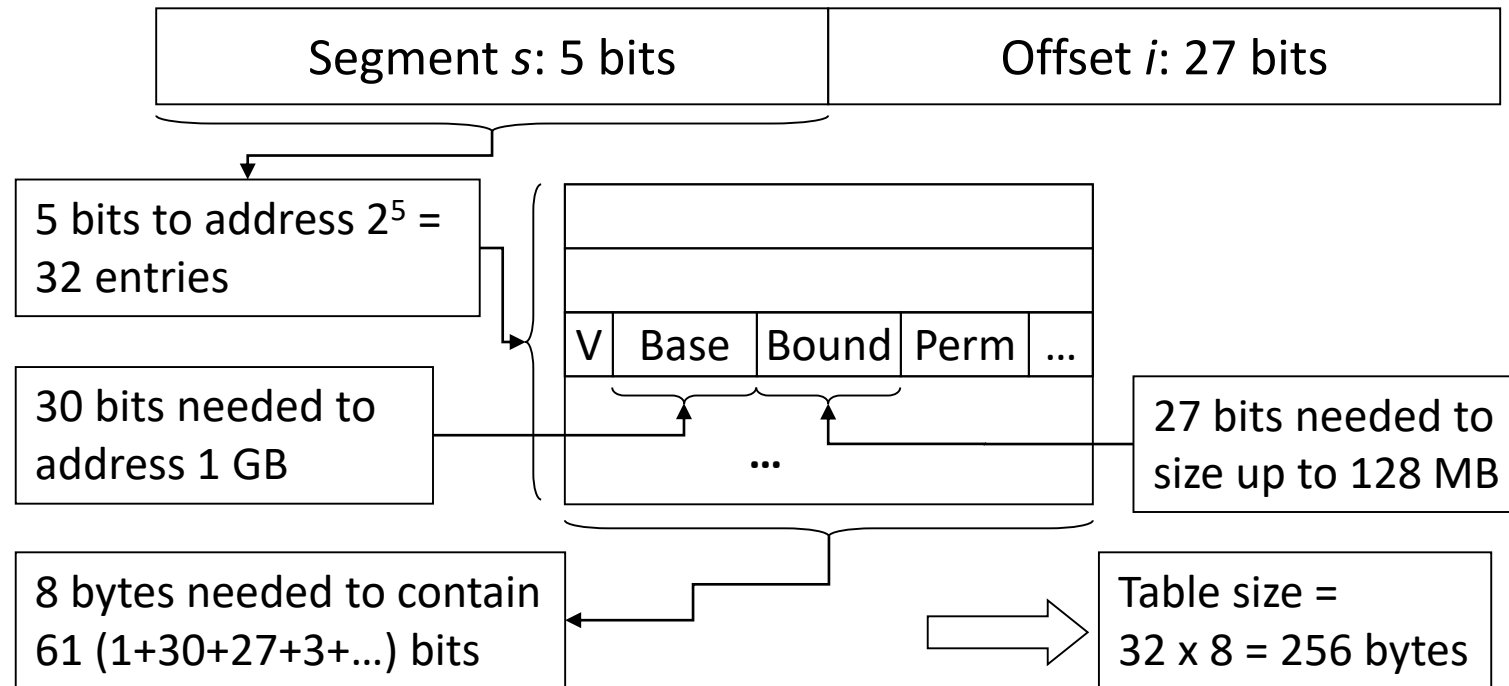
- Given 32 bit logical, 1 GB physical memory (max)
 - 5 bit segment number, 27 bit offset

Example of Sizing the Segment Table



- Given 32 bit logical, 1 GB physical memory (max)
 - 5 bit segment number, 27 bit offset

Example of Sizing the Segment Table




- Given 32 bit logical, 1 GB physical memory (max)
 - 5 bit segment number, 27 bit offset

Pros and Cons of Segmentation


- Pro: Each segment can be
 - located independently
 - separately protected
 - grown/shrunk independently
- Pro: Small segment table size
- Con: Variable-size allocation
 - Difficult to find holes in physical memory
 - External fragmentation

Defining Regions - Two Approaches

- Segmentation:
 - Partition address space and memory into segments
 - Segments have varying sizes



- Paging:
 - Partition address space and memory into pages
 - Pages are a constant, fixed size




Paging Vocabulary

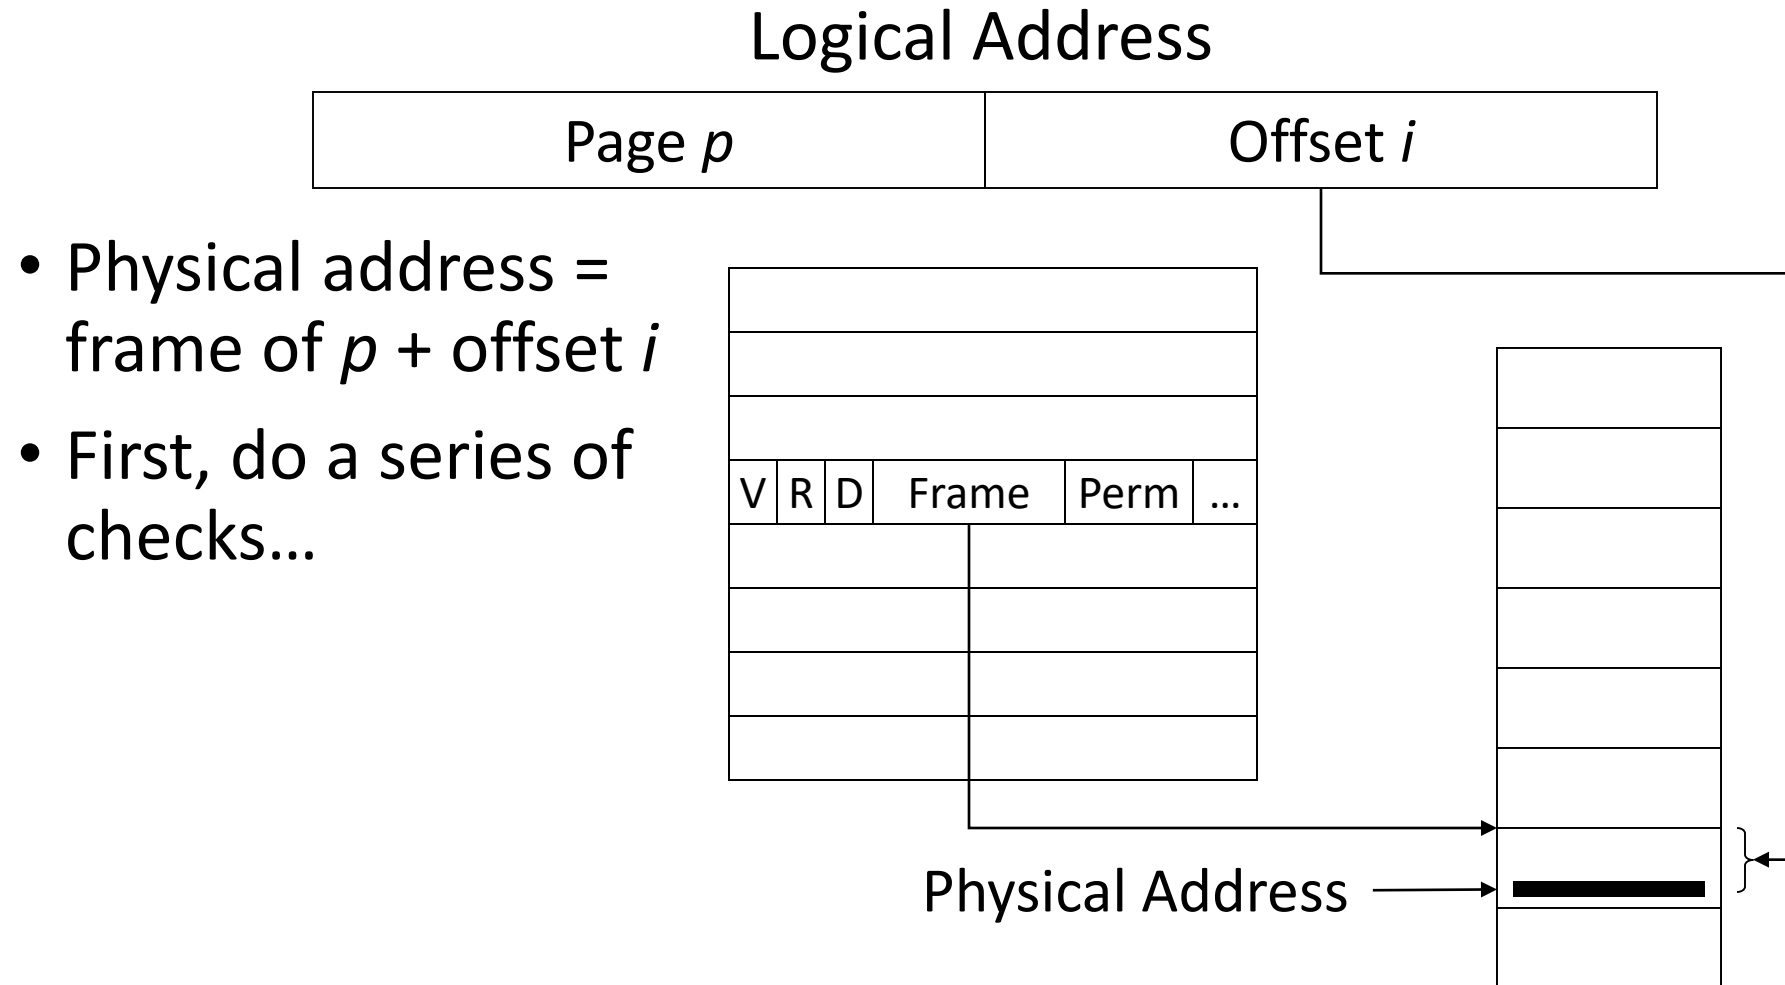
- For each process, the virtual address space is divided into fixed-size pages.
- For the system, the physical memory is divided into fixed-size frames.
- The size of a page is equal to that of a frame.
 - Often 4 KB in practice.

Page Table

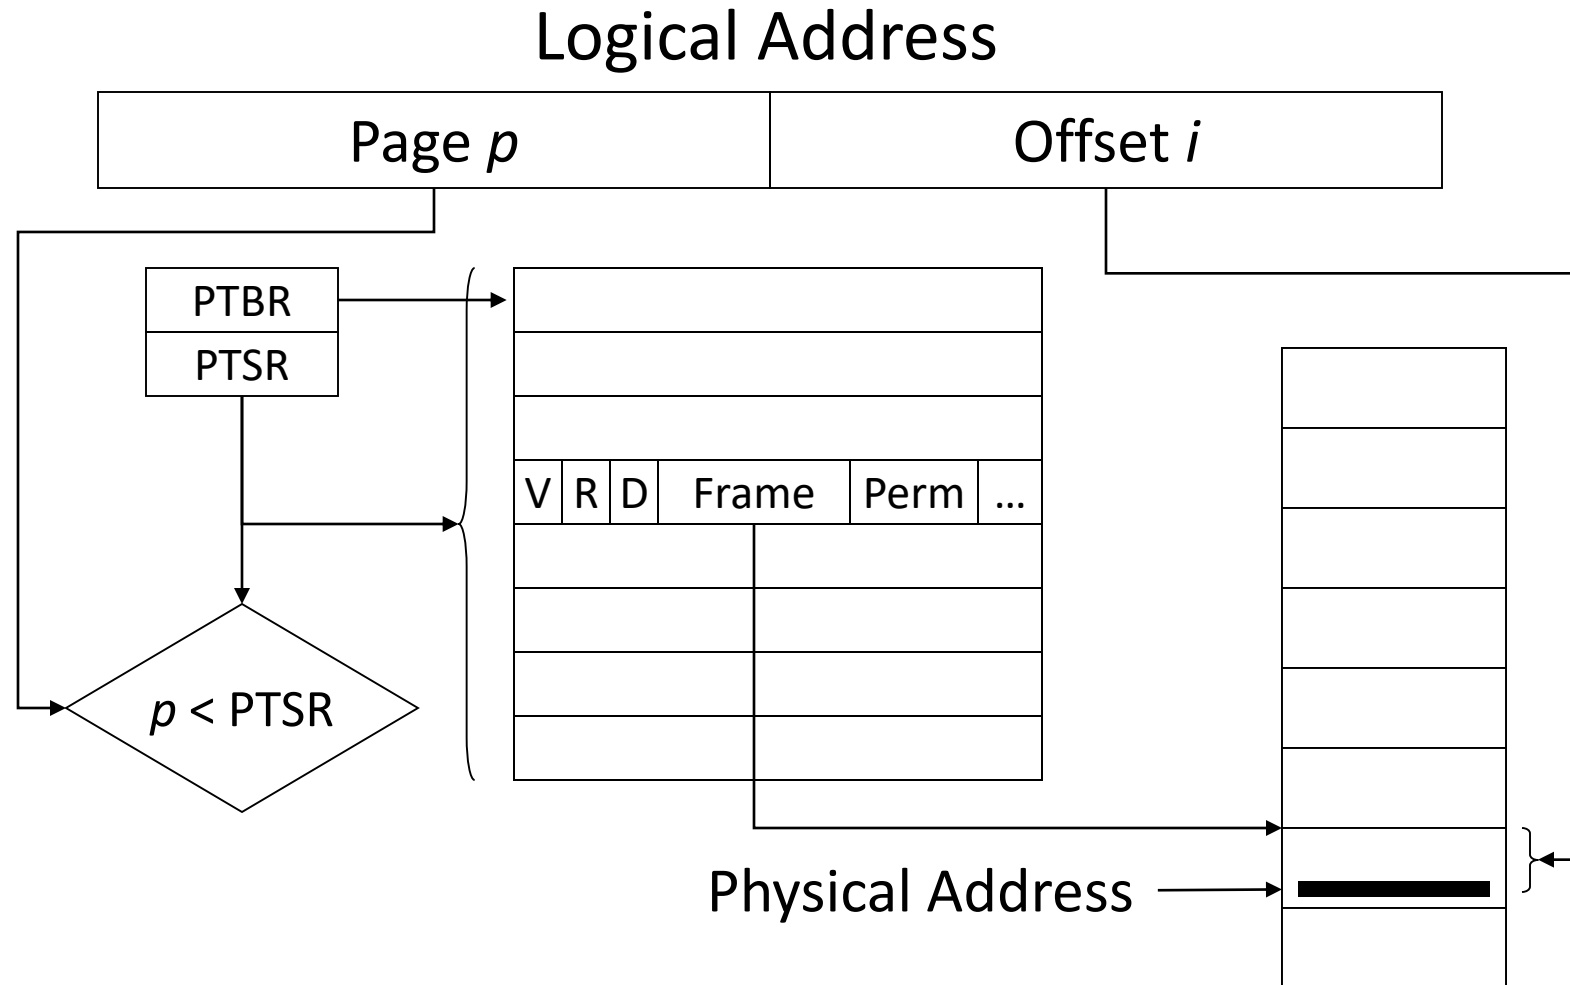
- One table per process
- Table parameters in memory
 - Page table base register
 - Page table size register
- Table entry elements
 - V: valid bit
 - R: referenced bit
 - D: dirty bit
 - Frame: location in phy mem
 - Perm: access permissions



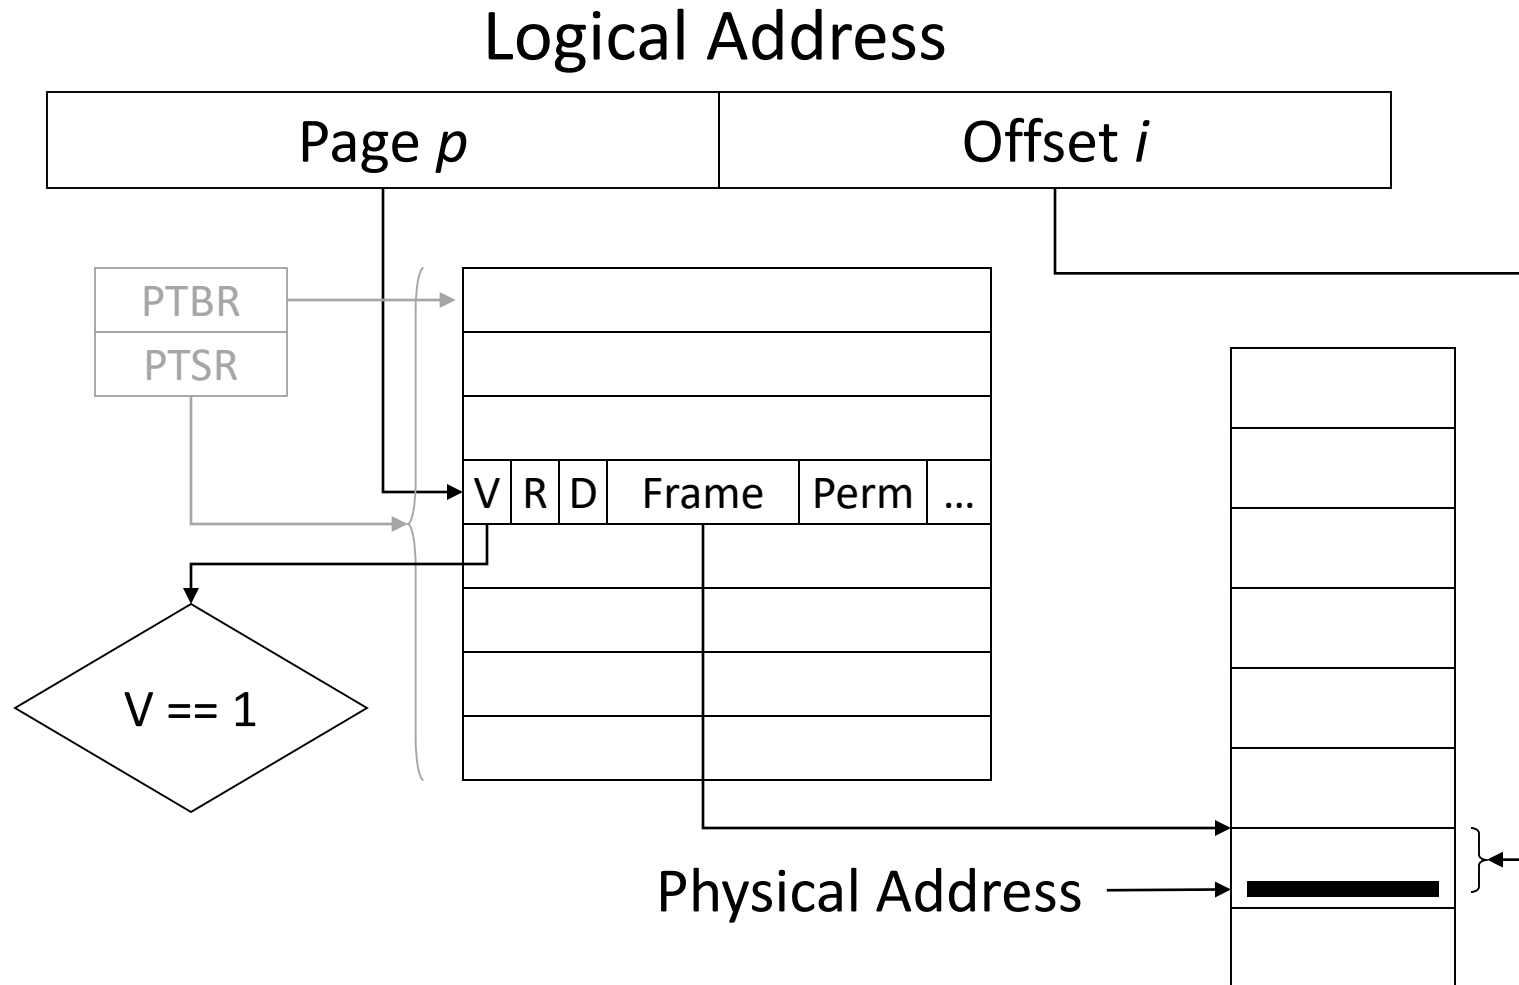
Address Translation



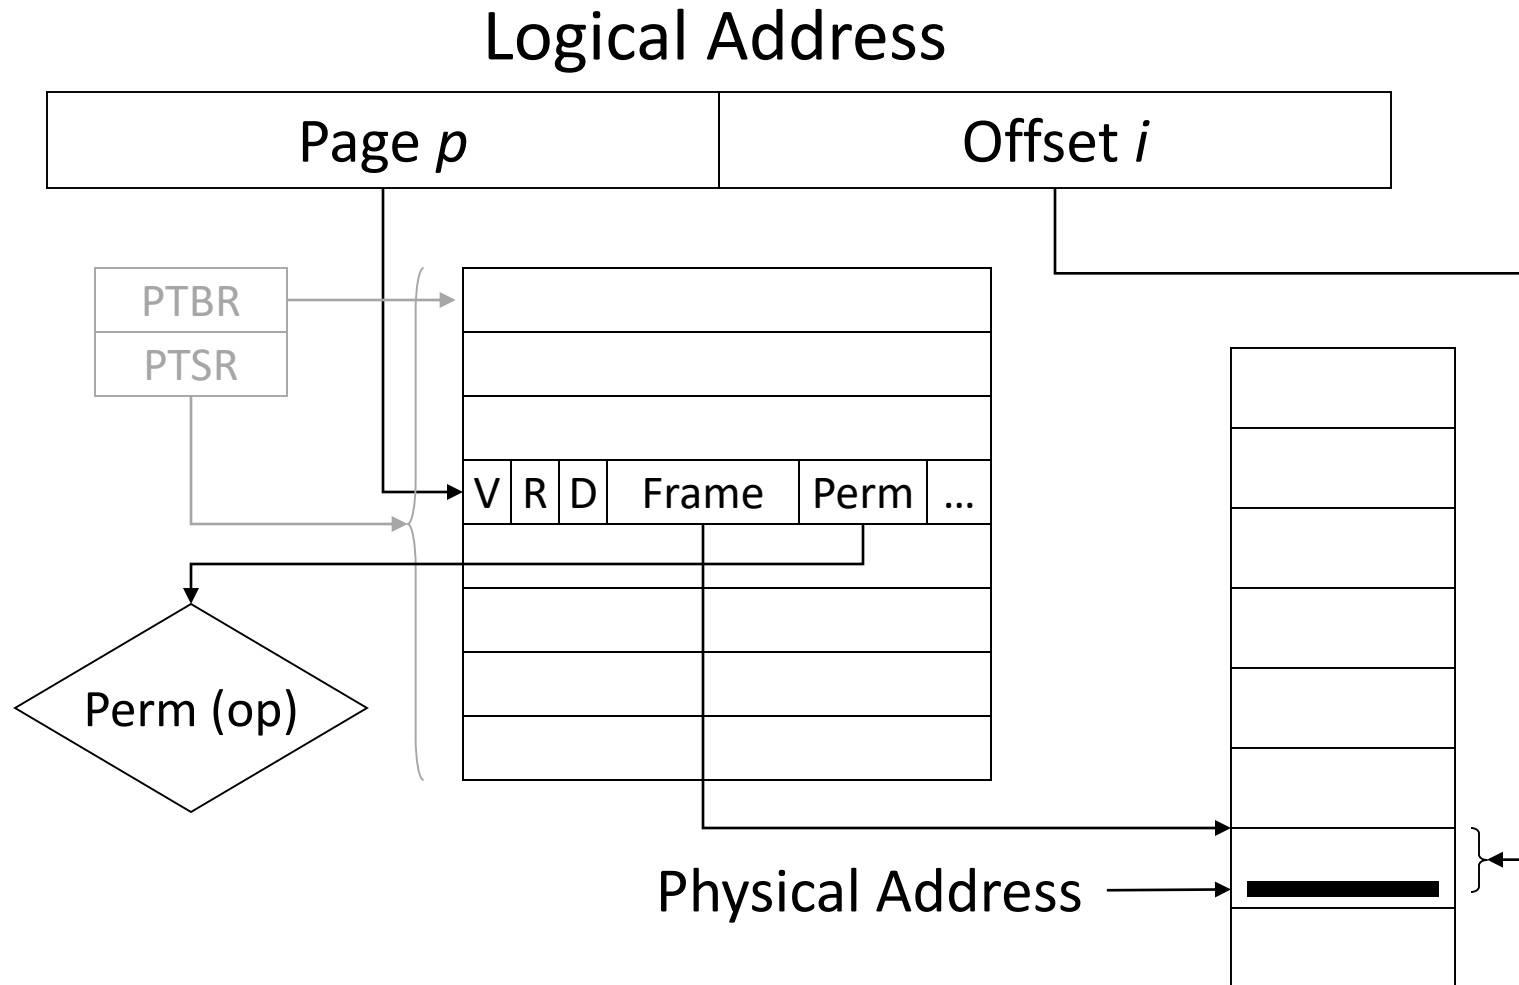
Check if Page p is Within Range




Check if Page Table Entry p is Valid



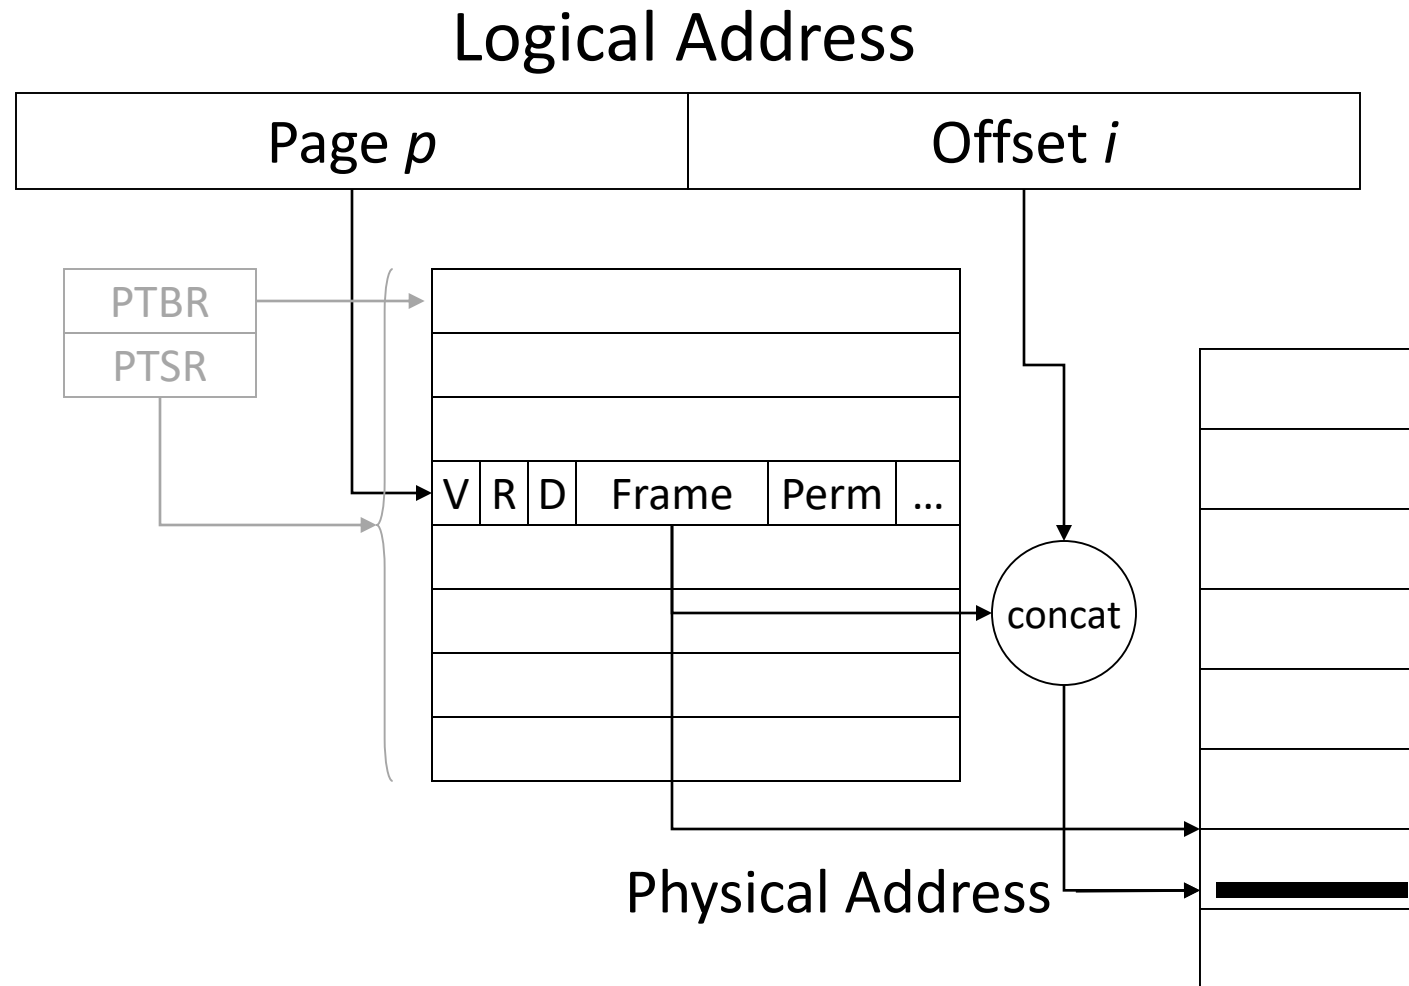
Check if Operation is Permitted



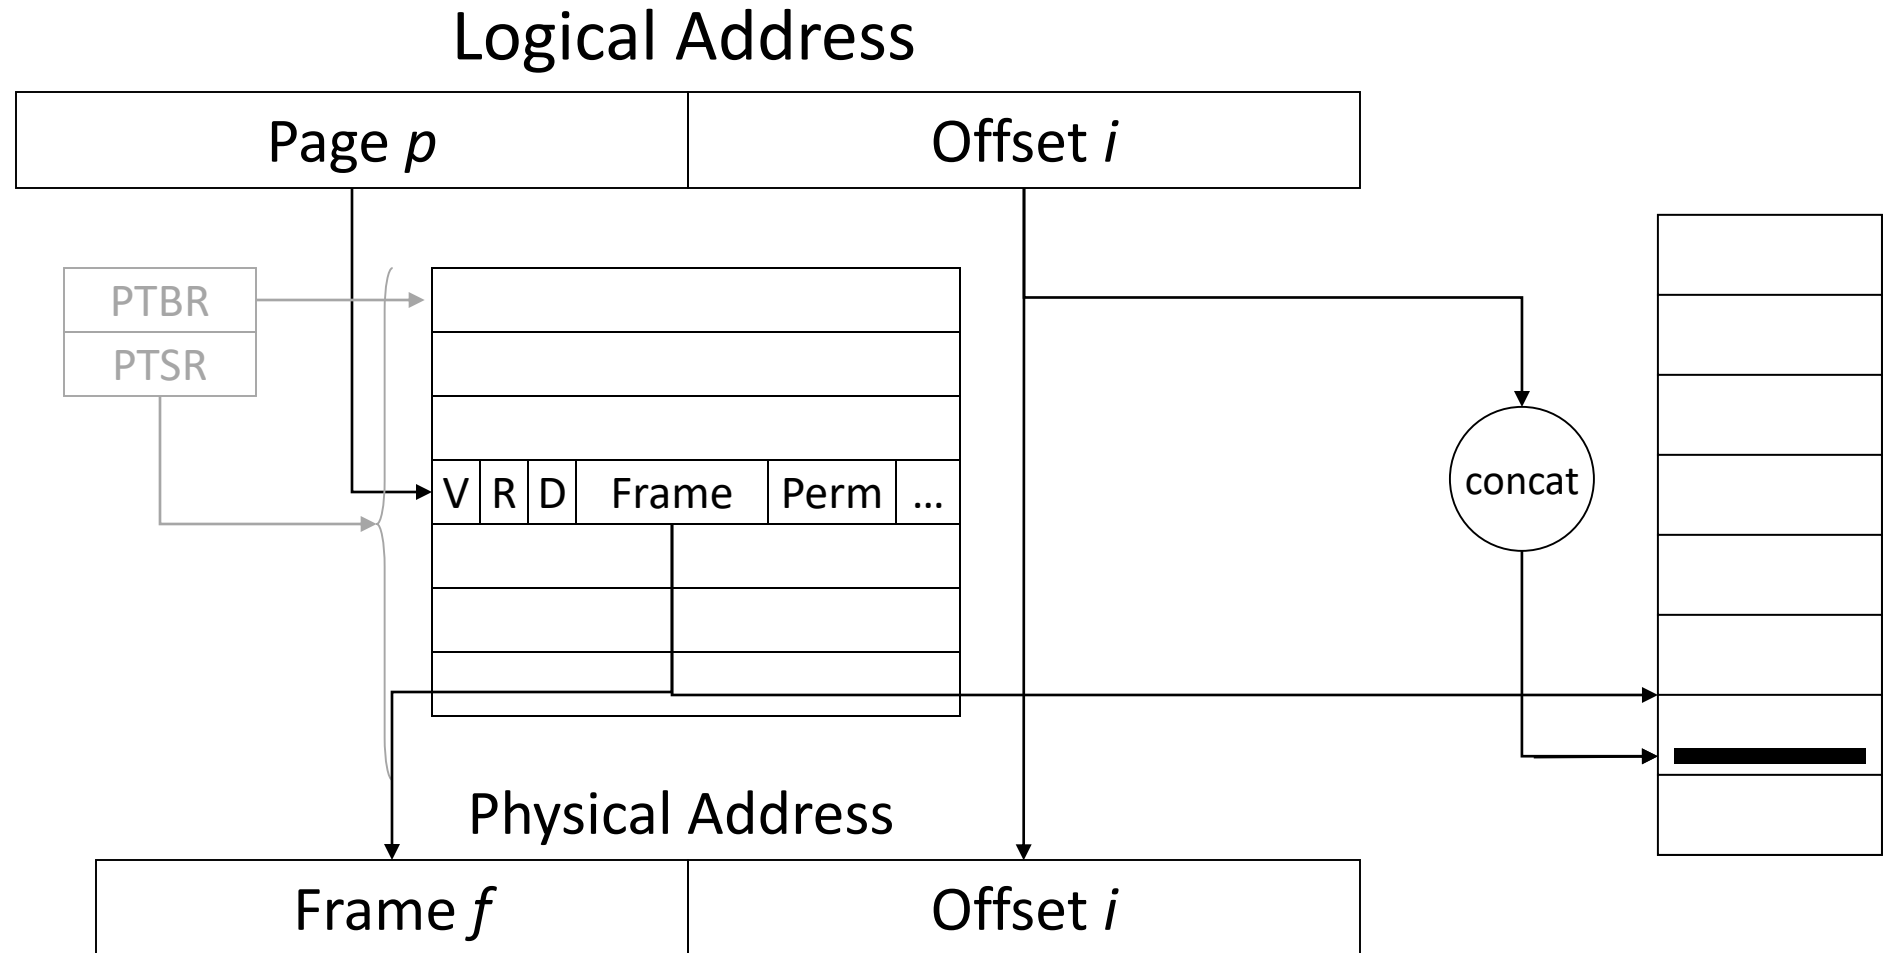
Translate Address



Translate Address

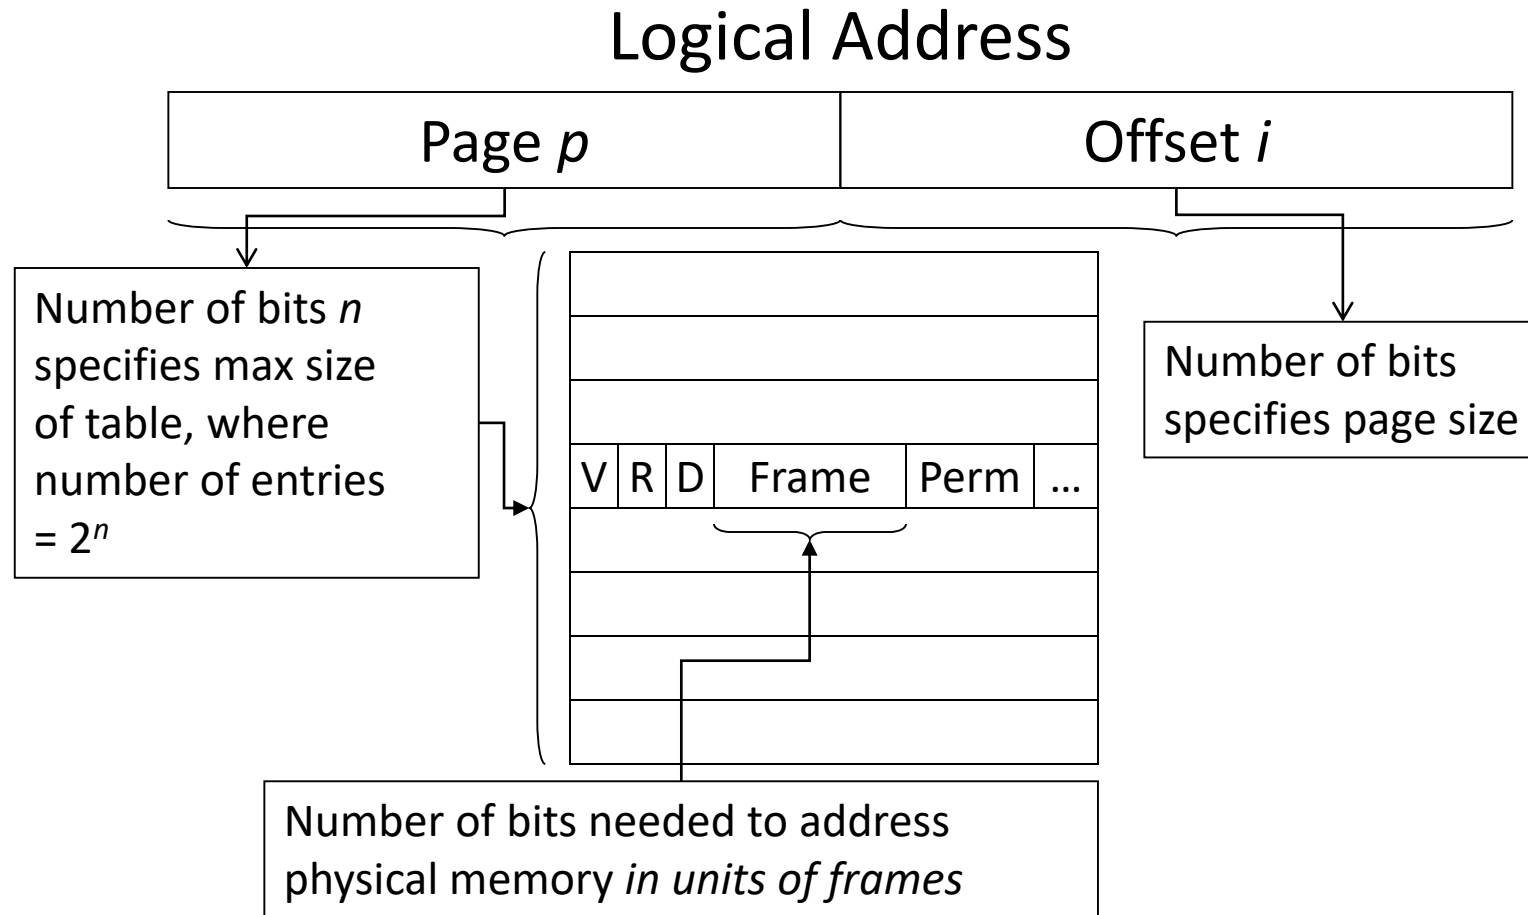


Physical Address by Concatenation




Frames are all the same size. Only need to store the *frame number* in the table, not exact address!

Sizing the Page Table

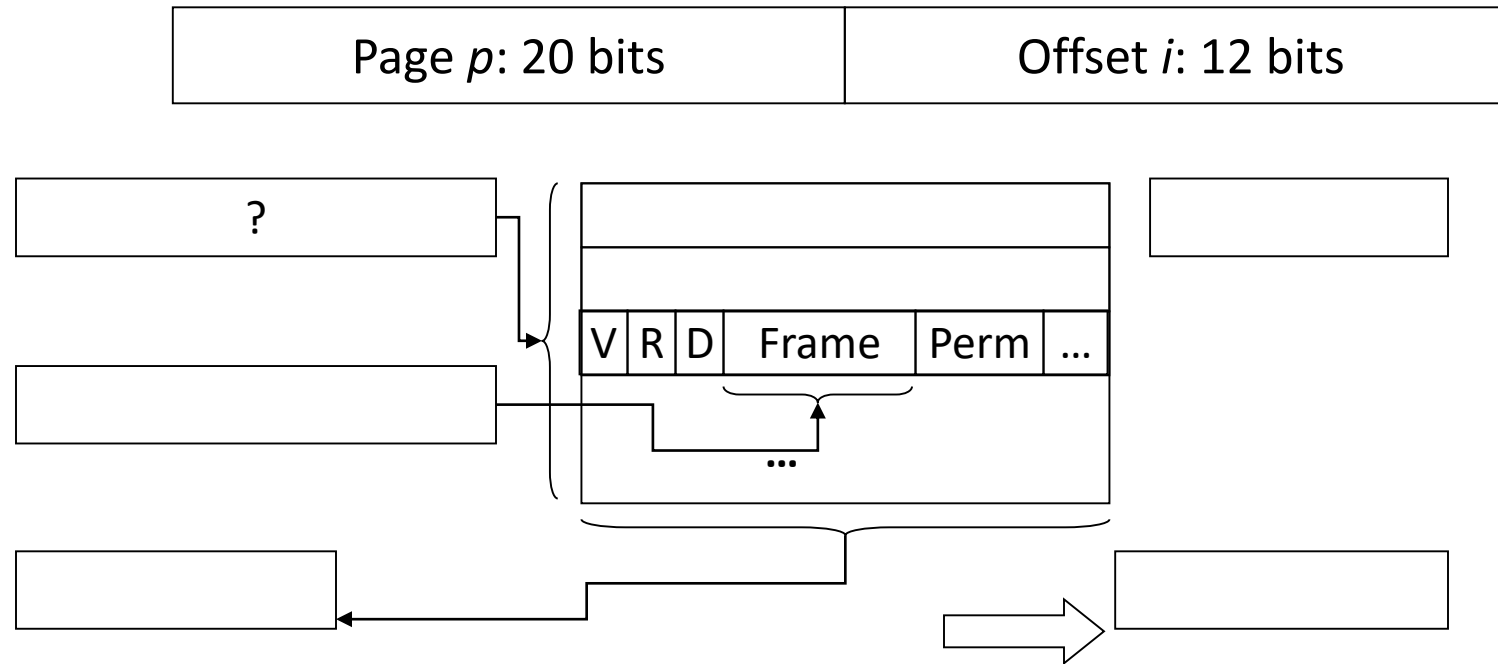


Example of Sizing the Page Table



- Given: 32 bit virtual addresses, 1 GB physical memory
 - Address partition: 20 bit page number, 12 bit offset

Example of Sizing the Page Table

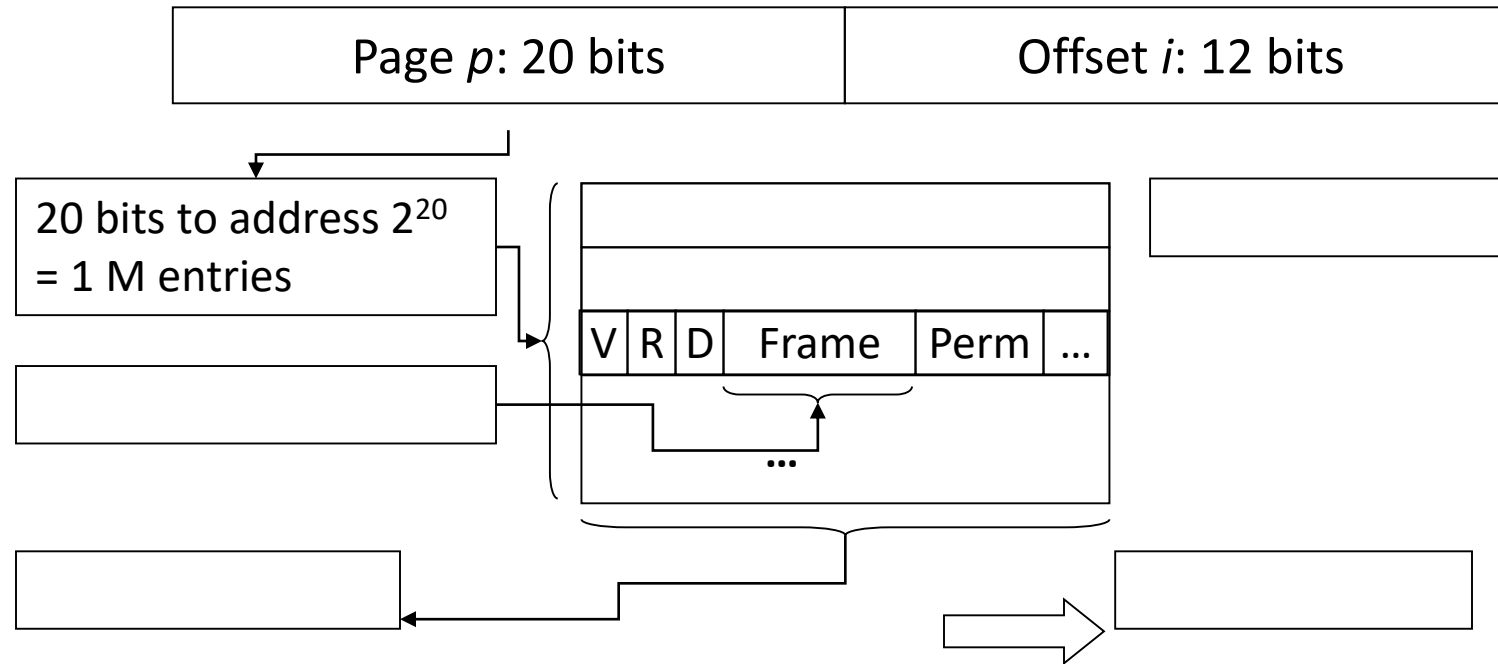


- Given: 32 bit virtual addresses, 1 GB physical memory
 - Address partition: 20 bit page number, 12 bit offset

How many entries (rows) will there be in this page table?

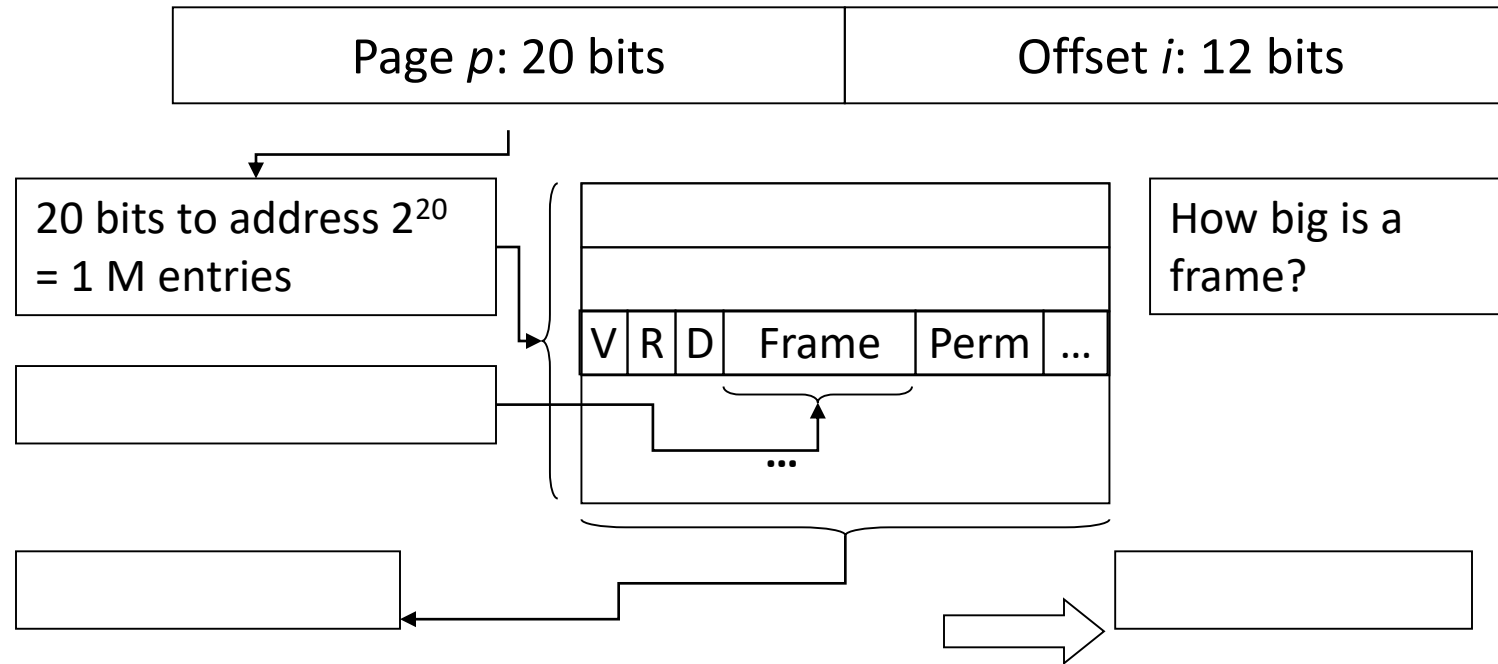
- A. 2^{12} , because that's how many the offset field can address
- B. 2^{20} , because that's how many the page field can address
- C. 2^{30} , because that's how many we need to address 1 GB
- D. 2^{32} , because that's the size of the entire address space

Example of Sizing the Page Table



- Given: 32 bit virtual addresses, 1 GB physical memory
 - Address partition: 20 bit page number, 12 bit offset

Example of Sizing the Page Table




- Given: 32 bit virtual addresses, 1 GB physical memory
 - Address partition: 20 bit page number, 12 bit offset

What will be the frame size, in bytes?

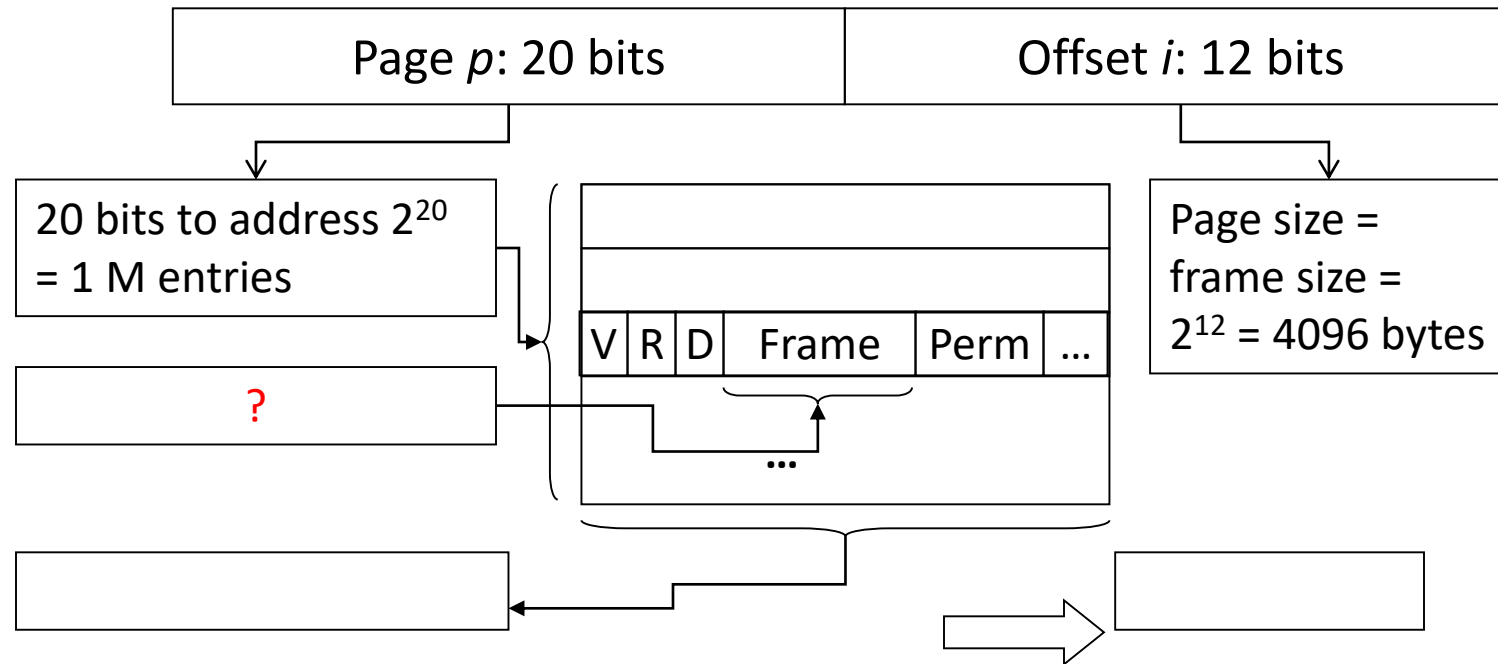
- A. 2^{12} , because that's how many bytes the offset field can address
- B. 2^{20} , because that's how many bytes the page field can address
- C. 2^{30} , because that's how many bytes we need to address 1 GB
- D. 2^{32} , because that's the size of the entire address space

Example of Sizing the Page Table




- Given: 32 bit virtual addresses, 1 GB physical memory
 - Address partition: 20 bit page number, 12 bit offset

How many bits do we need to store the frame number?



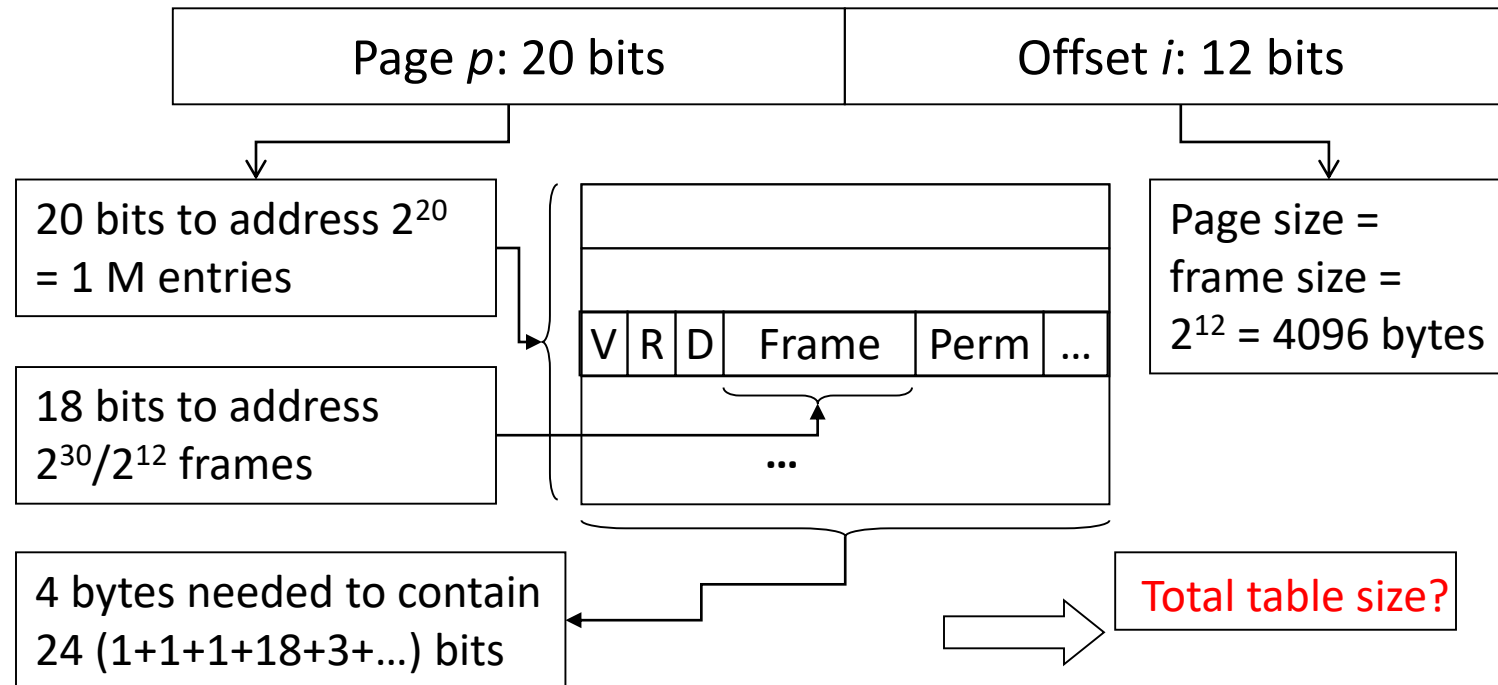
- Given: 32 bit virtual addresses, 1 GB physical memory
 - Address partition: 20 bit page number, 12 bit offset
- A: 12 B: 18 C: 20 D: 30 E: 32

Example of Sizing the Page Table



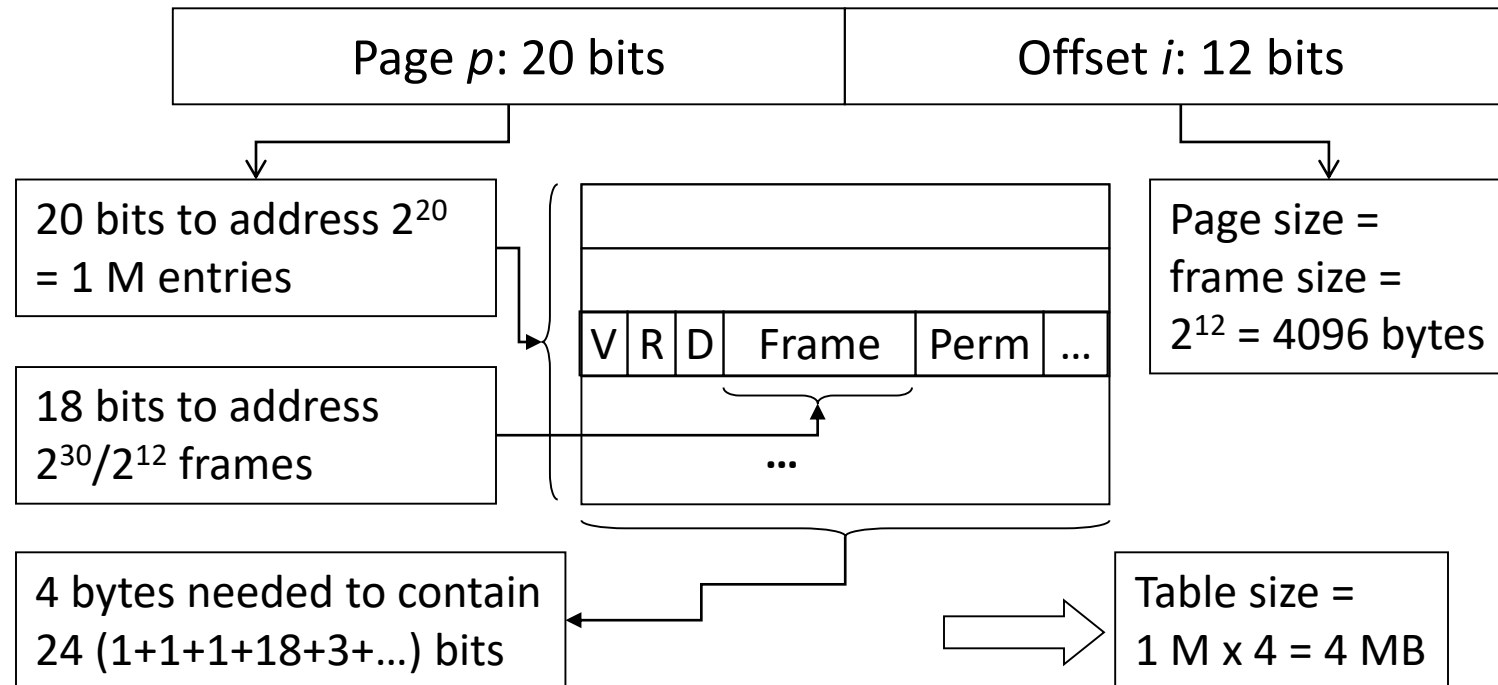
- Given: 32 bit virtual addresses, 1 GB physical memory
 - Address partition: 20 bit page number, 12 bit offset

Example of Sizing the Page Table



- Given: 32 bit virtual addresses, 1 GB physical memory
 - Address partition: 20 bit page number, 12 bit offset

Example of Sizing the Page Table




- 4 MB of bookkeeping for *every process*?
 - 200 processes -> 800 MB just to store page tables...

Pros and Cons of Paging

- Pro: Fixed-size pages and frames
 - No external fragmentation
 - No difficult placement decisions
- Con: large table size
- Con: *maybe* internal fragmentation

x86: Hybrid Approach

- Design:
 - Multiple lookups: first in segment table, which points to a page table.
 - Extra level of indirection.
- Reality:
 - All segments are max physical memory size
 - Segments effectively unused, available for “legacy” reasons.
 - (Mostly) disappeared in x86-64



Outstanding Problems

- Mostly considering paging from here on.
1. Page tables are way too big. Most processes don't need that many pages, can't justify a huge table.
 2. Adding indirection hurts performance.

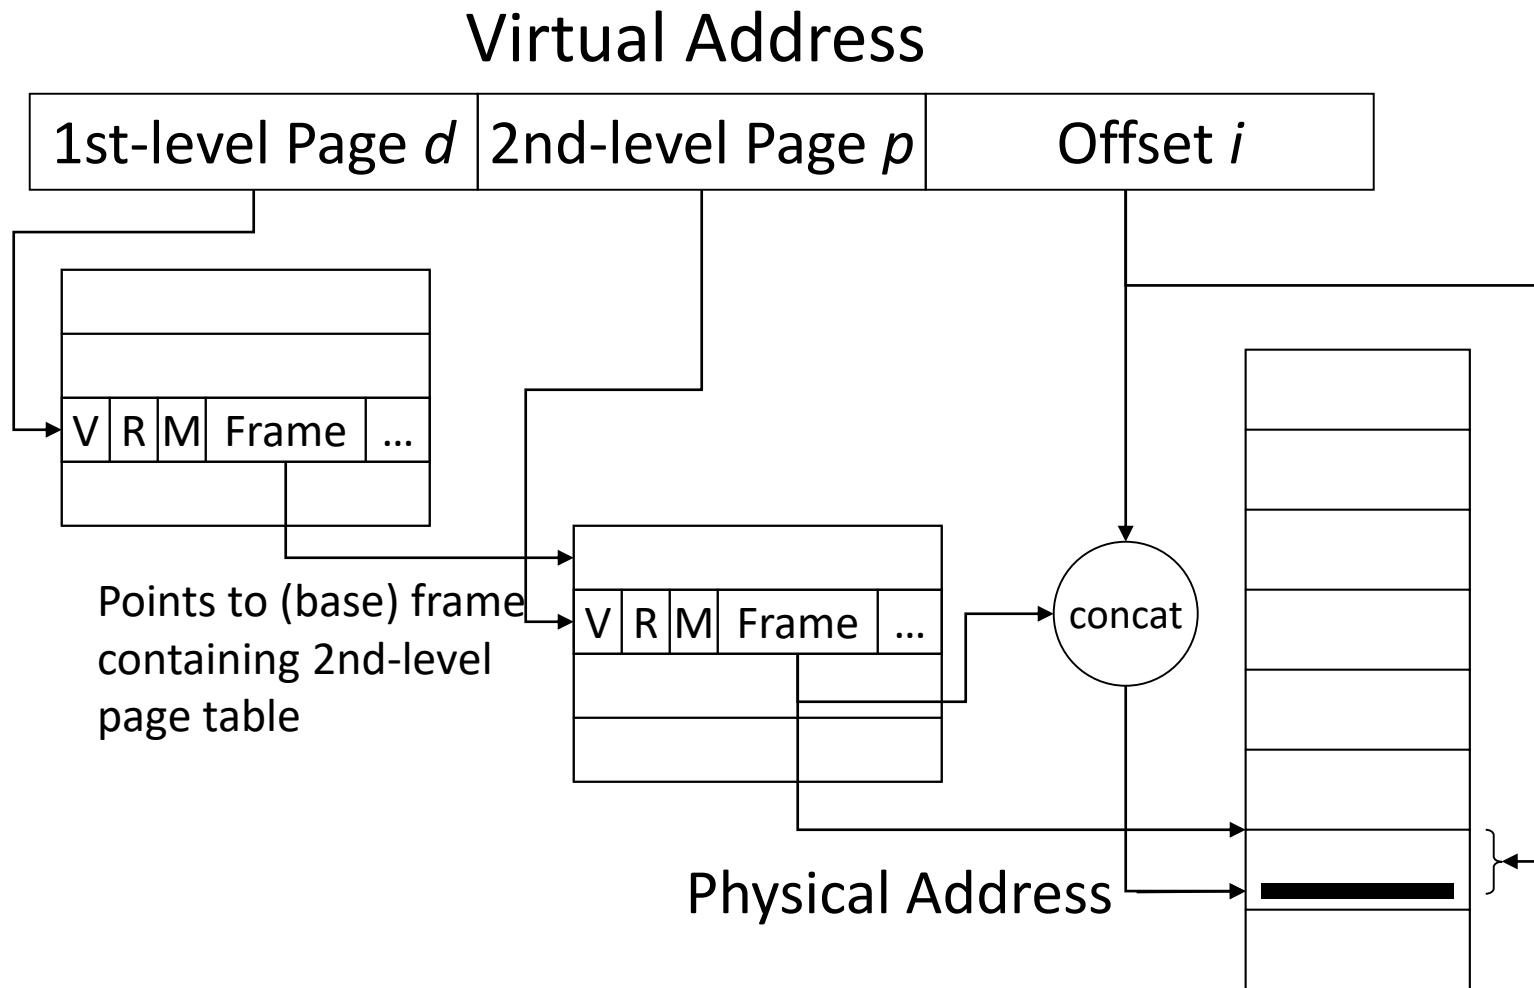
Outstanding Problems

- Mostly considering paging from here on.
1. Page tables are way too big. Most processes don't need that many pages, can't justify a huge table.
 2. Adding indirection hurts performance.



Solution:
MORE indirection!

Multi-Level Page Tables


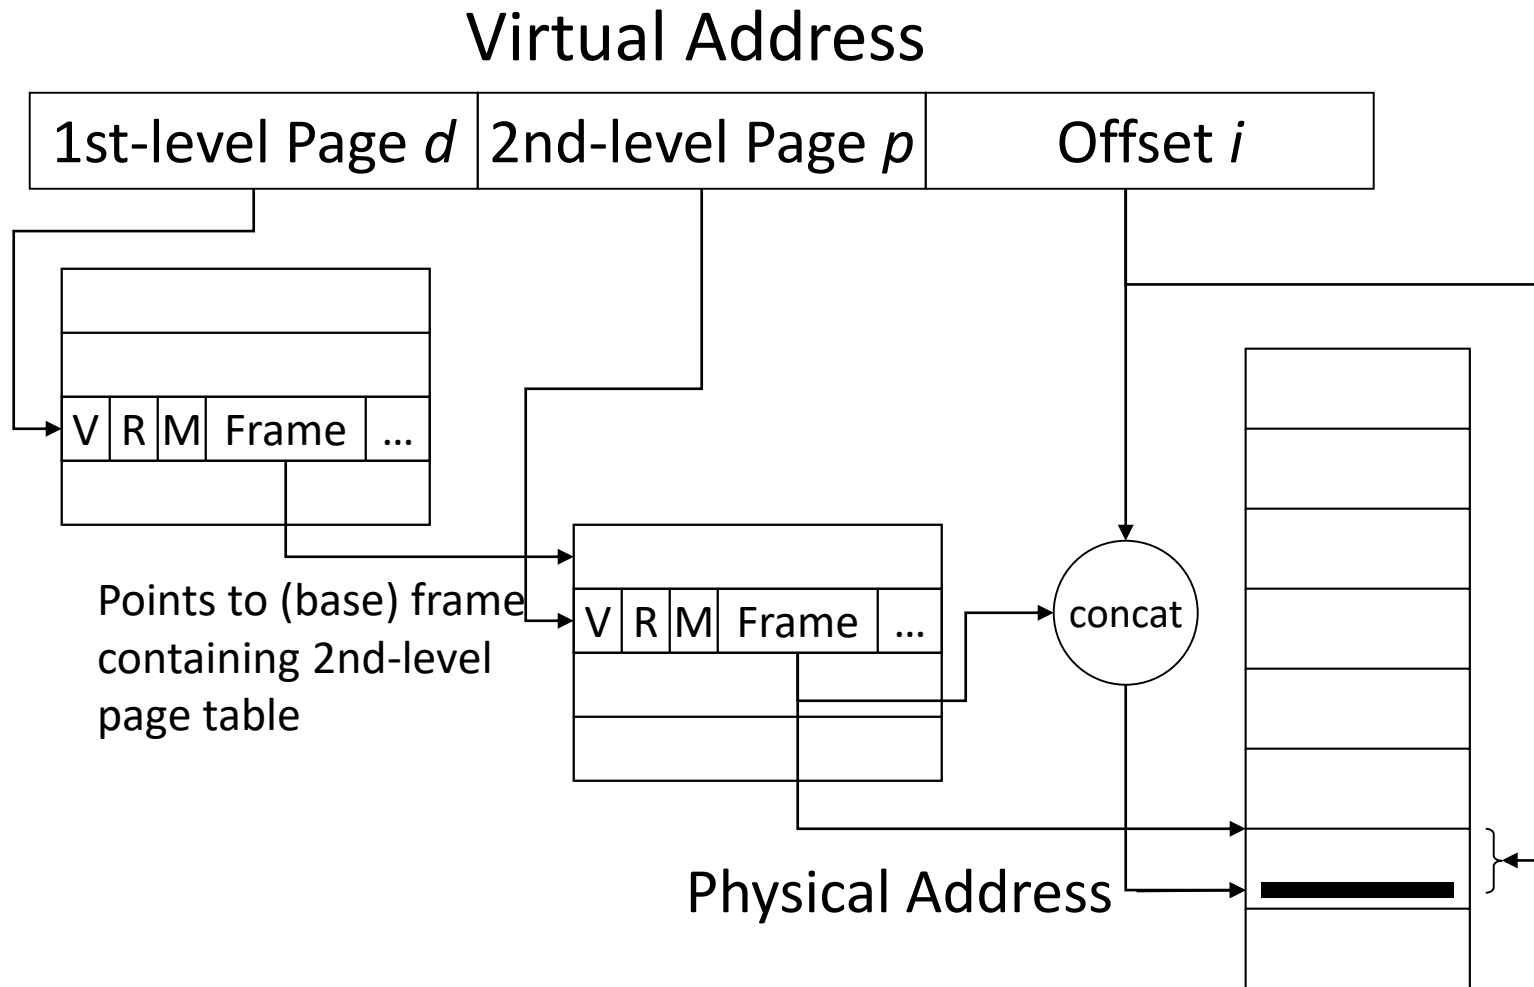


Insight: VAS is typically sparsely populated.

Idea: every process gets a page directory (1st-level table)

Only allocate 2nd-level tables when the process is using that VAS region!

Multi-Level Page Tables



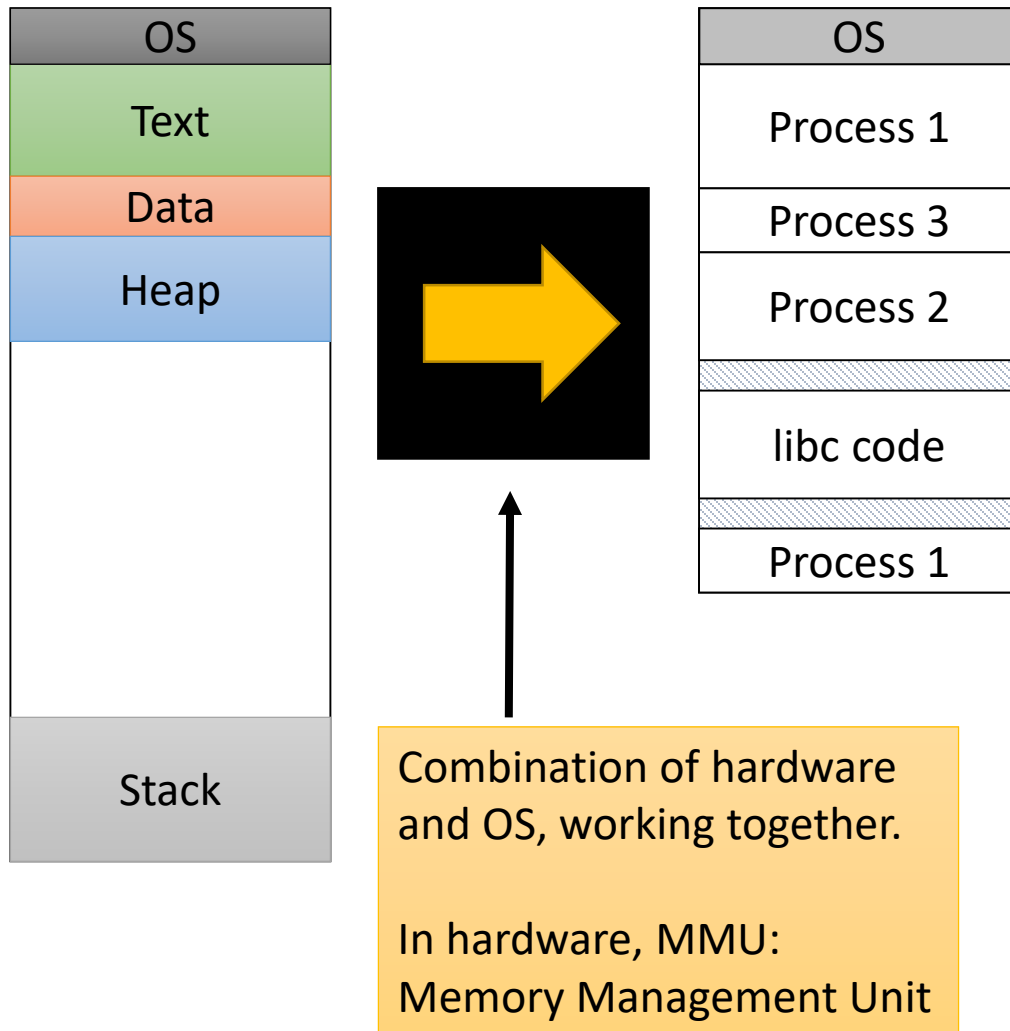
Multi-Level Page Tables

- With only a single level, the page table must be large enough for the largest processes.
- Extra level of indirection:
 - WORSE performance – more memory accesses
 - Much better memory efficiency – process's page table is proportional to how much of the VAS it's using.
- Small process -> low page table storage
- Large process -> high page table storage, needed it anyway

Translation Cost

- Each application memory access now requires multiple accesses!
- Suppose memory takes 100 ns to access.
 - one-level paging: 200 ns
 - two-level paging: 300 ns
- Solution: Add hardware, take advantage of locality...
 - Most references are to a small number of pages
 - Keep translations of these in high-speed memory


Memory Management Unit



- When a process tries to use memory, send the address to MMU.
- MMU will do as much work as it can. If it knows the answer, great!
- If it doesn't, trigger exception (OS gets control), consult software table.

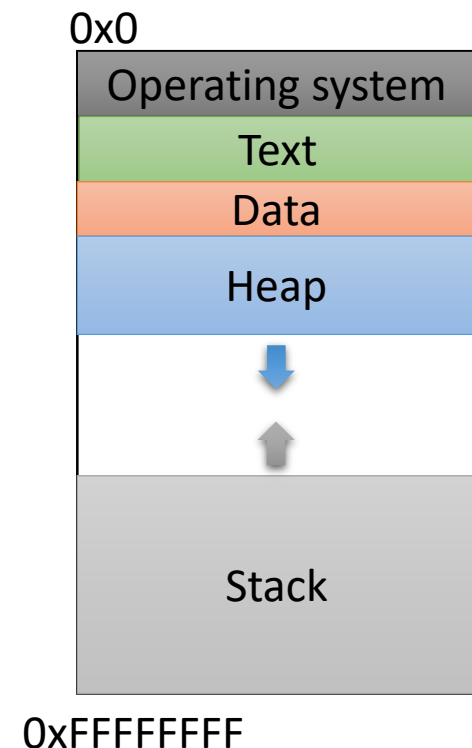
Translation Look-aside Buffer (TLB)

- Fast memory mapping cache inside MMU keeps most recent translations
 - If key matches, get frame number quickly
 - otherwise, wait for normal translation (in parallel)



Recall: Context Switching Performance

- Even though it's fast, context switching is expensive:
 1. time spent is 100% overhead
 2. must invalidate other processes' resources (caches, memory mappings)
 3. kernel must execute – it must be accessible in memory
- Also recall: Advantage of threads
 - Threads all share one process VAS



Translation Cost with TLB

- Cost is determined by
 - Speed of memory: ~ 100 nsec
 - Speed of TLB: ~ 10 nsec
 - Hit ratio: fraction of refs satisfied by TLB, $\sim 95\%$
- Speed to access memory with no address translation: 100 nsec
- Speed to access memory with address translation:
 - TLB miss: 300 nsec (200% slowdown)
 - TLB hit: 110 nsec (10% slowdown)
 - Average: $110 \times 0.95 + 300 \times 0.05 = 119.5$ nsec

TLB Design Issues

- The larger the TLB...
 - the higher the hit rate
 - the slower the response
 - the greater the expense
 - the larger the space (in MMU, on chip)
- TLB has a major effect on performance!
 - Must be flushed on context switches
 - Alternative: tagging entries with PIDs

Summary

- Many options for translation mechanism: segmentation, paging, hybrid, multi-level paging. All of them: level(s) of *indirection*.
- Simplicity of paging makes it most common today.
- Multi-level page tables improve memory efficiency – page table bookkeeping scales with process VAS usage.
- TLB in hardware MMU exploits locality to improve performance